

Solving Problems by Search:

Problem Formulation and Uninformed Search

CZ3005: Artificial Intelligence

Shen Zhiqi



Outline

- ❑ How to define a problem-solving agent?
- ❑ Types of problem
- ❑ Problem formulation
- ❑ Informed search strategies

Problem-Solving Agent

Rational goal-based agent:

- Performance measure defined in terms of satisfying goals

Goal formulation

- Define and organize objectives (goal states)

Problem formulation

- Define what states and actions (transitions) to consider

Search for a solution

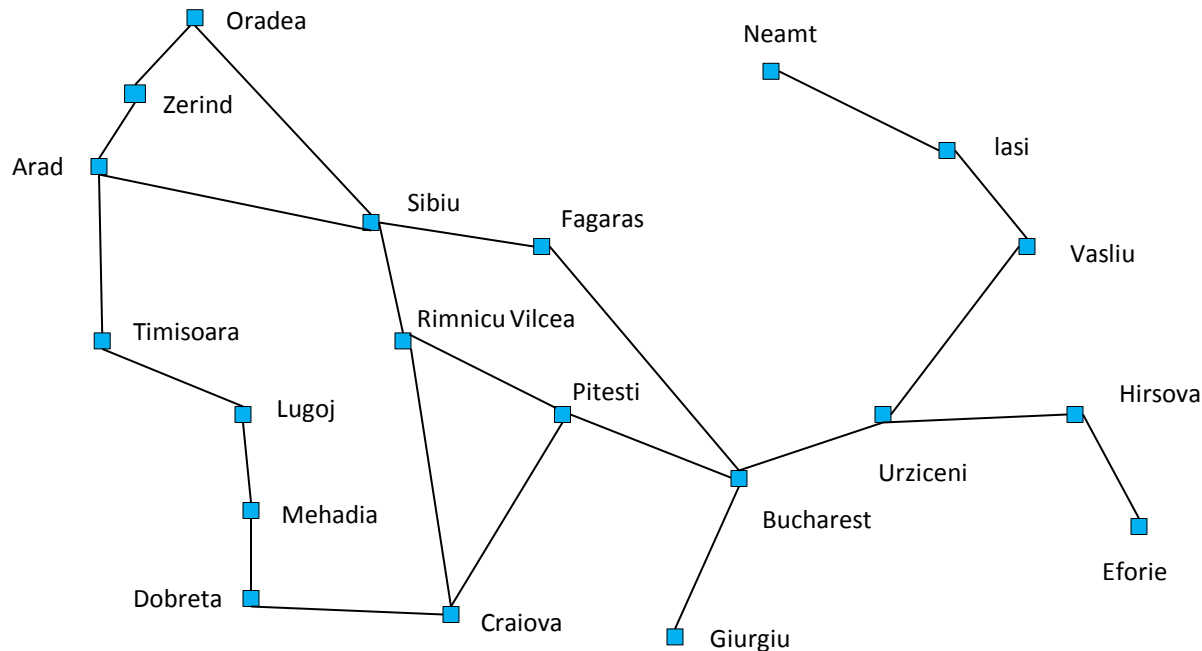
- Find a sequence of actions that lead to a goal state

Execution

- Actually carry out the recommended actions

Goal-based Agent: Example

On holiday in Romania



- ❑ Currently in Arad (**Initial state**). Flight leaves tomorrow from Bucharest.
- ❑ **Goal**: be in Bucharest (other factors: cost, time, most scenic route, etc)
- ❑ **State**: be in a city (defined by the map)
- ❑ **Action**: transition between states (highways defined by the map)

Design of Problem-Solving Agent

Idea:

- ❑ Systematically considers the **expected outcomes** of different possible sequences of actions that lead to states of known value
- ❑ Choose the best one
 - shortest journey from A to B?
 - most cost effective journey from A to B?

Steps:

1. Goal formulation
2. Problem formulation
3. Search process

No knowledge → uninformed search

Knowledge → informed search

4. Action execution (follow the recommended route)

Algorithm

Function SIMPLE-PROBLEM-SOLVING-AGENT(*p*) **returns** an action

Inputs: *p*, a percept

Static: *s*, an action sequence, initially empty

state, some description of the current world state

g, a goal, initially null

problem, a problem formulation

state \leftarrow UPDATE-STATE(*state*, *p*)

if *s* is empty **then**

g \leftarrow FORMULATE-GOAL(*state*)

problem \leftarrow FORMULATE-PROBLEM(*state*, *g*)

s \leftarrow SEARCH(*problem*)

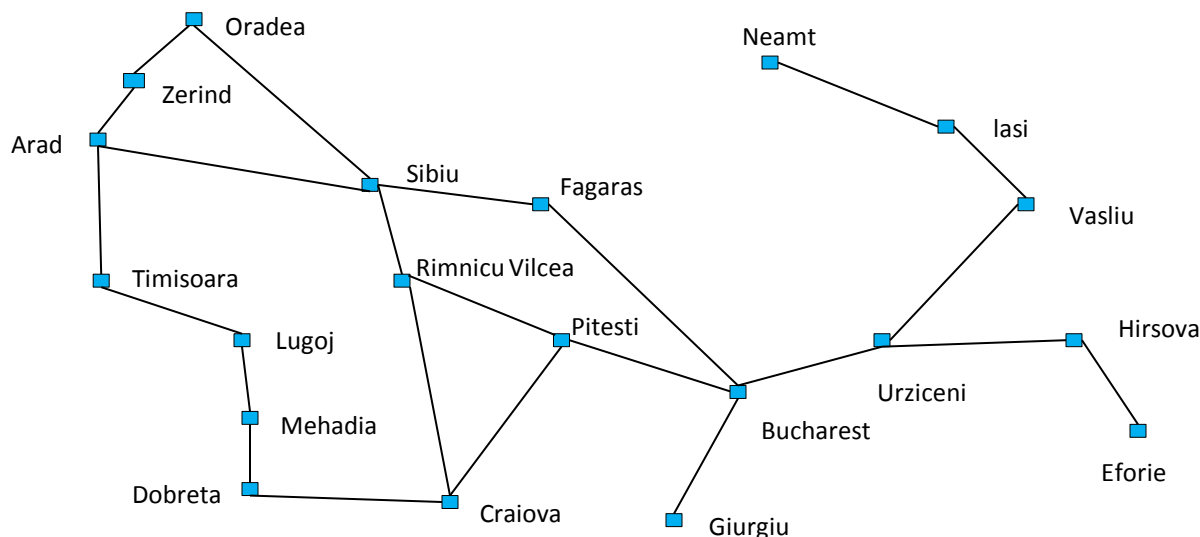
action \leftarrow RECOMMENDATION(*s*, *state*)

s \leftarrow REMAINDER(*s*, *state*)

return *action*

Example: Romania

- ❑ **Goal:** be in Bucharest
- ❑ Formulate problem:
 - ❑ **states:** various cities
 - ❑ **actions:** drive between cities
- ❑ **Solution:**
 - ❑ sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest




Example: Vacuum Cleaner Agent



- ❑ Robotic vacuum cleaners move autonomously
- ❑ Some can come back to a docking station to charge their batteries
- ❑ A few are able to empty their dust containers into the dock as well

Example: A Simple Vacuum World

- Two locations, each location may or may not contain dirt, and the agent may be in one location or the other



- 8 possible world states
- Possible actions: left, right, and suck
- Goal: clean up all dirt → Two goal states, i.e. {7, 8}


Problem Types

The **problem type** depends on how much knowledge the agent has concerning its actions and the state that it is in.

- ❑ **Single-State** Problem
- ❑ **Multiple-State** Problem
- ❑ **Contingency** Problem

Single-State Problem


- ❑ **Accessible** world state (sensory information is available)
- ❑ **Known** outcome of action (deterministic)



- ❑ e.g.: start in #5
 - ❑ Solution: **right, suck**

Multiple-State Problem


- ❑ **Inaccessible** world state (with limited sensory information):
agent only knows which **sets** of states it is in
- ❑ **Known** outcome of action (deterministic)



- ❑ e.g.: start in {1, 2, 3, 4, 5, 6, 7, 8}
 - ❑ Action **right** goes to {2, 4, 6, 8}
 - ❑ Solution: **right, suck, left, suck**

Contingency Problem

- Limited or no sensory information (**inaccessible**)
- Limited agent knowledge, action result is **not predictable**
- Effect of action depends on what is found to be true through perception/monitoring (**non-deterministic**)
 - Suppose action **suck** **may** deposits dirt at a clean carpet



- e.g. start in #4, **suck** → reach {2,4}, solution: **left**, suck
- e.g. start in {1, 3} and have only local dirt sensor ,solution:
 - **suck, right**
 - **suck** only if there is dirt there
- problem solving requires sensing **during** the execution phase
 - e.g., keep your eyes open while walking(contingency)

Well-Defined Formulation

Definition of a problem:

- The information used by an agent to decide what to do

Specification:

- Initial state
- Action set, i.e. available actions (successor functions)
- State space, i.e. states reachable from the initial state
 - Solution **path**: sequence of actions from one state to another
- Goal test predicate
 - Single state, enumerated list of states, abstract properties
- Cost function
 - Path cost $g(n)$, sum of all (action) step costs along the path

Solution:

- A path (a sequence of operators leading) from the **Initial-State** to a state that satisfies the **Goal-Test**

Measuring Problem-Solving Performance

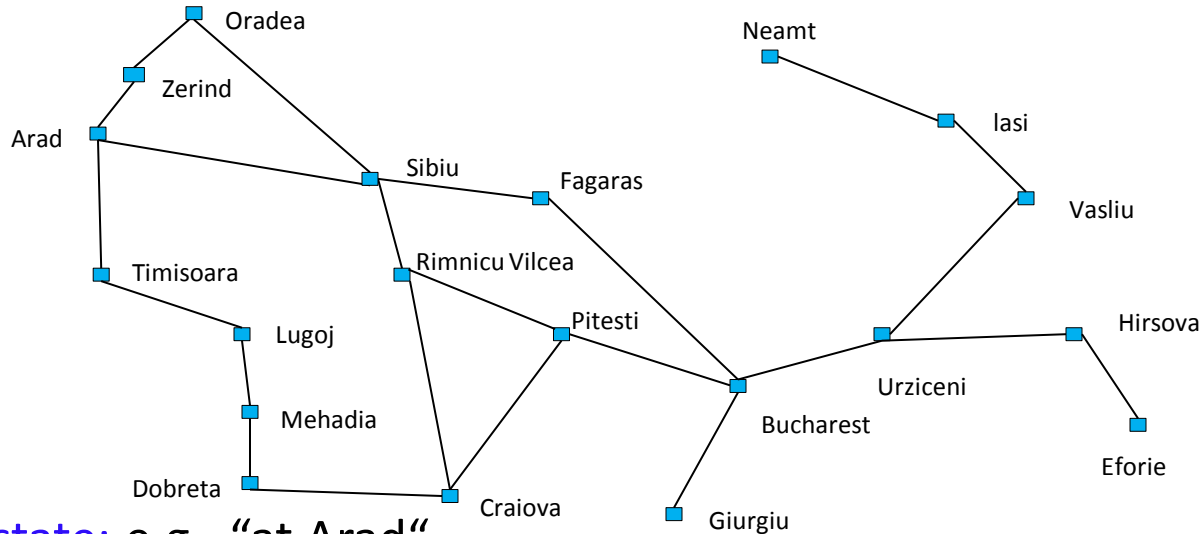
Search Cost:

- What does it cost to find the solution?
 - e.g. How long (time)? How many resources used (memory)?

Total cost of problem-solving

- Search cost ("offline") + Execution cost ("online")
- Trade-offs often required
 - Search a very long time for the optimal solution, or
 - Search a shorter time for a "good enough" solution


Single-State Problem Example



- ❑ Initial state: e.g., “at Arad”
- ❑ Set of possible actions and the corresponding next states
 - ❑ e.g., Arad → Zerind
- ❑ Goal test:
 - ❑ explicit (e.g., x = “at Bucharest”)
- ❑ Path cost function
 - ❑ e.g., sum of distances, number of operators executed solution: a sequence of operators leading from the initial state to a goal state

Example: Vacuum World (Single-state Version)

- ❑ **Initial state:** one of the eight states shown previously
- ❑ **Actions:** left, right, suck
- ❑ **Goal test:** no dirt in any square
- ❑ **Path cost:** 1 per action



Multiple-State Problem Formulation

- ❑ Initial state set
- ❑ Set of possible actions and the corresponding sets of next states
- ❑ Goal test
- ❑ Path cost function

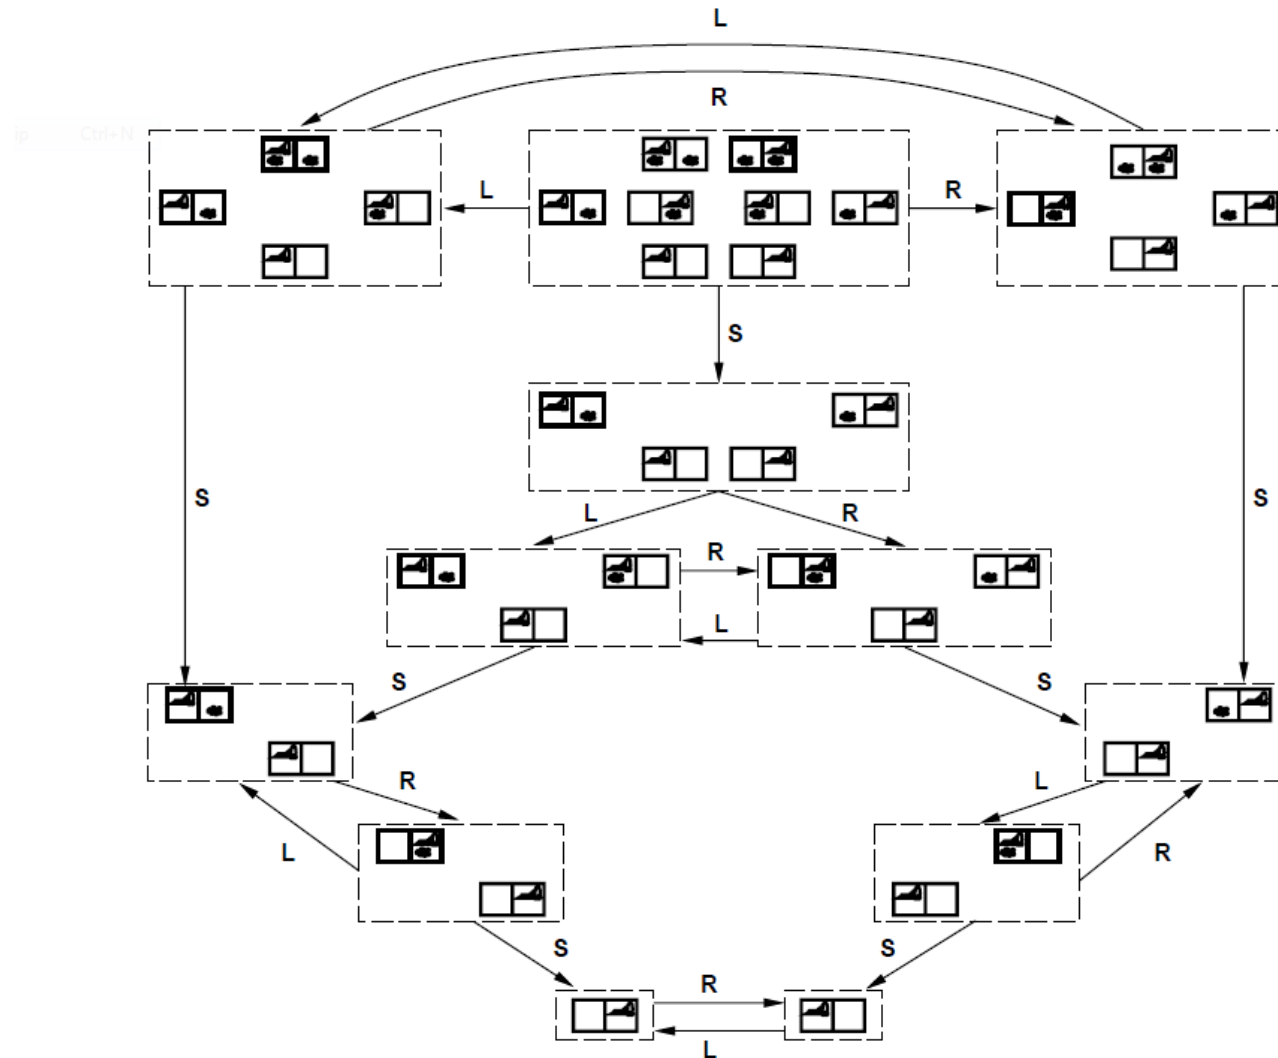
Solution:

- ❑ a path (connecting sets of states) that leads to a set of states all of which are goal states

Example: Vacuum World (Multiple-state Version)

- ❑ **States**: subset of the eight states
- ❑ **Operators**: left, right, suck
- ❑ **Goal test**: all states in state set have no dirt
- ❑ **Path cost**: 1 per operator

Example: Vacuum World (Multiple-state Version)



Example: 8-puzzle

5	4	
6	1	8
7	3	2


Start state

1	2	3
8		4
7	6	5

Goal state

- ❑ **States**: integer locations of tiles
 - ❑ number of states = $9!$
- ❑ **Actions**: move blank left, right, up, down
- ❑ **Goal test**: = goal state (given)
- ❑ **Path cost**: 1 per move

Example: 8-queens



- ❑ **States:** Any arrangement of 0 to 8 queens on the board
- ❑ **Actions:** Add a queen to any empty square
- ❑ **Goal test:** 8 queens are on the board, none attacked
- ❑ **Path cost:** Not necessary

Real-World Problems

Route finding problems:

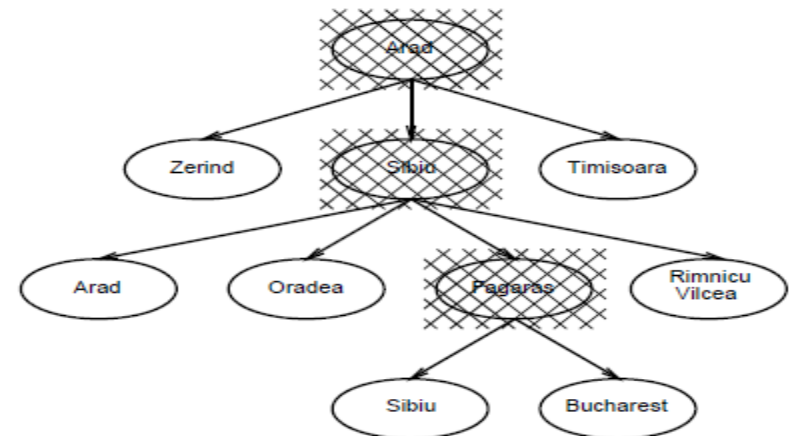


- ❑ Routing in computer networks
- ❑ Robot navigation
- ❑ Automated travel advisory
- ❑ Airline travel planning

Touring problems:

- ❑ Traveling Salesperson problem
- ❑ “Shortest tour”: visit every city exactly once

Search Algorithms

- Exploration of state space by generating successors of already-explored states
- **Expanding** the states
 - Frontier: candidate nodes for expansion
 - Explored set



Search Algorithms...

```
Function GENERAL-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
      else expand the node and add the resulting nodes to the search tree
  end
```

Search Strategies

A **strategy** is defined by picking the order of node expansion

Strategies are evaluated along the following dimensions:

- ❑ **completeness**
 - ❑ does it always find a solution if one exists?
- ❑ **time complexity**
 - ❑ how long does it take to find a solution: the number of nodes generated
- ❑ **space complexity**
 - ❑ maximum number of nodes in memory
- ❑ **optimality**
 - ❑ does it always find the best (least-cost) solution?

Branching factor

- ❑ Maximum number of successors of any node
- ❑ Or average branching factor (Tutorial Q1(c))

Uninformed vs Informed

Uninformed search strategies


- ❑ use **only** the information available in the problem definition
- ❑ Breadth-first search
- ❑ Uniform-cost search
- ❑ Depth-first search
- ❑ Depth-limited search
- ❑ Iterative deepening search

Informed search strategies

- ❑ use **problem-specific knowledge** to guide the search
- ❑ usually more efficient

Breadth-First Search

Expand **shallowest** unexpanded node which can be implemented by a First-In-First-Out (FIFO) queue



Denote

- b : maximum branching factor of the search tree
- d : depth of the least-cost solution

Complete: Yes

Optimal: Yes when **all** step costs equally

Complexity of BFS

Branching factor


- Hypothetical state-space, where every node can be expanded into b new nodes, solution of path-length d
- Time: $1 + b + b^2 + b^3 + \dots + b^d = O(b^d)$
- Space: (keeps every node in memory) $O(b^d)$ are equal

Depth	Nodes	Time	Memory
0	1	1 millisecond	100 bytes
2	111	.1 seconds	11 kilobytes
4	11111	11 seconds	1 megabyte
6	10^6	18 minutes	111 megabytes
8	10^8	31 hours	11 gigabytes
10	10^{10}	128 days	1 terabyte
12	10^{12}	35 years	111 terabytes
14	10^{14}	3500 years	11111 terabytes

Uniform-Cost Search

To consider **edge costs**, expand unexpanded node with the **least** path cost g

- ❑ Modification of breath-first search
- ❑ Instead of First-In-First-Out (FIFO) queue, using a priority queue with path cost $g(n)$ to order the elements
- ❑ BFS = UCS with $g(n) = \text{Depth}(n)$



Uniform-Cost Search...

- ❑ Complete: Yes
- ❑ Time: # of nodes with path cost $g \leq$ cost of optimal solution (eqv. # of nodes pop out from the priority queue)
- ❑ Space: # of nodes with path cost $g \leq$ cost of optimal solution
- ❑ Optimal: Yes

Depth-First Search

Denote

- m : maximum depth of the state space

Complete:

- infinite-depth spaces: No
- finite-depth spaces with loops: No
 - with repeated-state checking: Yes
- finite-depth spaces without loops: Yes

Time: $O(b^m)$

- if solutions are dense, may be much faster than breadth-first

Space: $O(bm)$

Optimal: No


Depth-Limited Search

To avoid infinite searching, Depth-first search with a **cutoff** on the max depth / of a path

- ❑ Complete: Yes, if $I \geq d$
- ❑ Time: $O(b^I)$
- ❑ Space: $O(bI)$
- ❑ Optimal: No

Iterative Deepening Search

Iteratively estimate the max depth / of DLS one-by-one



Iterative Deepening Search...

```
Function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution sequence
  inputs: problem, a problem

  for depth 0 to  $\infty$  do
    if DEPTH-LIMITED-SEARCH(problem, depth) succeeds then return its result
  end
  return failure
```

- ❑ Complete: Yes
- ❑ Time: $O(b^d)$
- ❑ Space: $O(bd)$
- ❑ Optimal: Yes

Summary (we make assumptions for optimality)

Criterion	Breadth-first	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional(if applicable)
Time	b^d	b^d	b^m	b^l	b^d	$b^{d/2}$
Space	b^d	b^d	bm	bl	bd	$b^{d/2}$
Optimal	Yes	Yes	No	No	Yes	Yes
Complete	Yes	Yes	No	Yes, if $l \geq d$	Yes	Yes