




CE/CZ2005 and CPE205/CSC205: Operating Systems – Lab Experiment 3

Contact: Bingsheng He
(bshe@ntu.edu.sg)

Outline

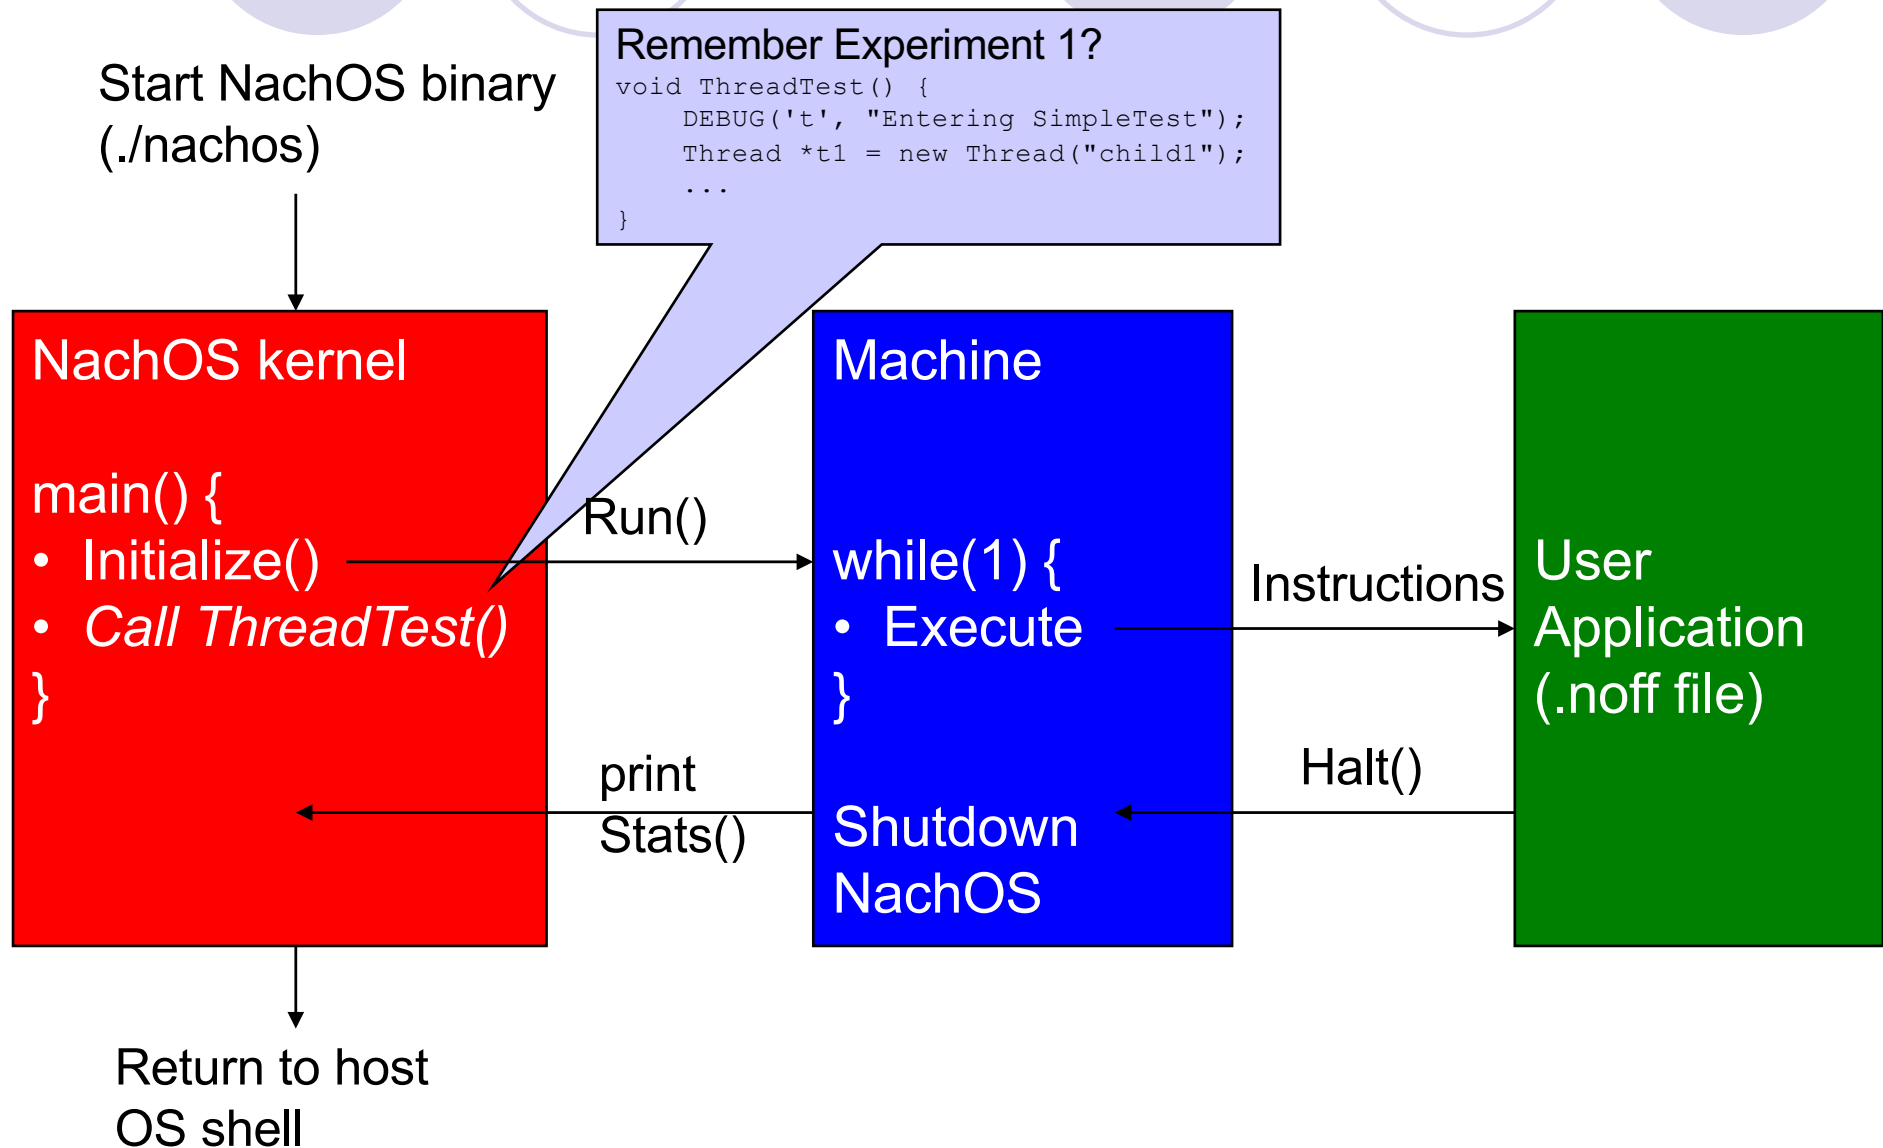



- Difference between Lab 1 and 2 to Lab 3
- Address translation in NachOS
- Discussion of Experiment 3



Difference between Labs 1 and 2 to Lab 3

NachOS – Remember how it works?

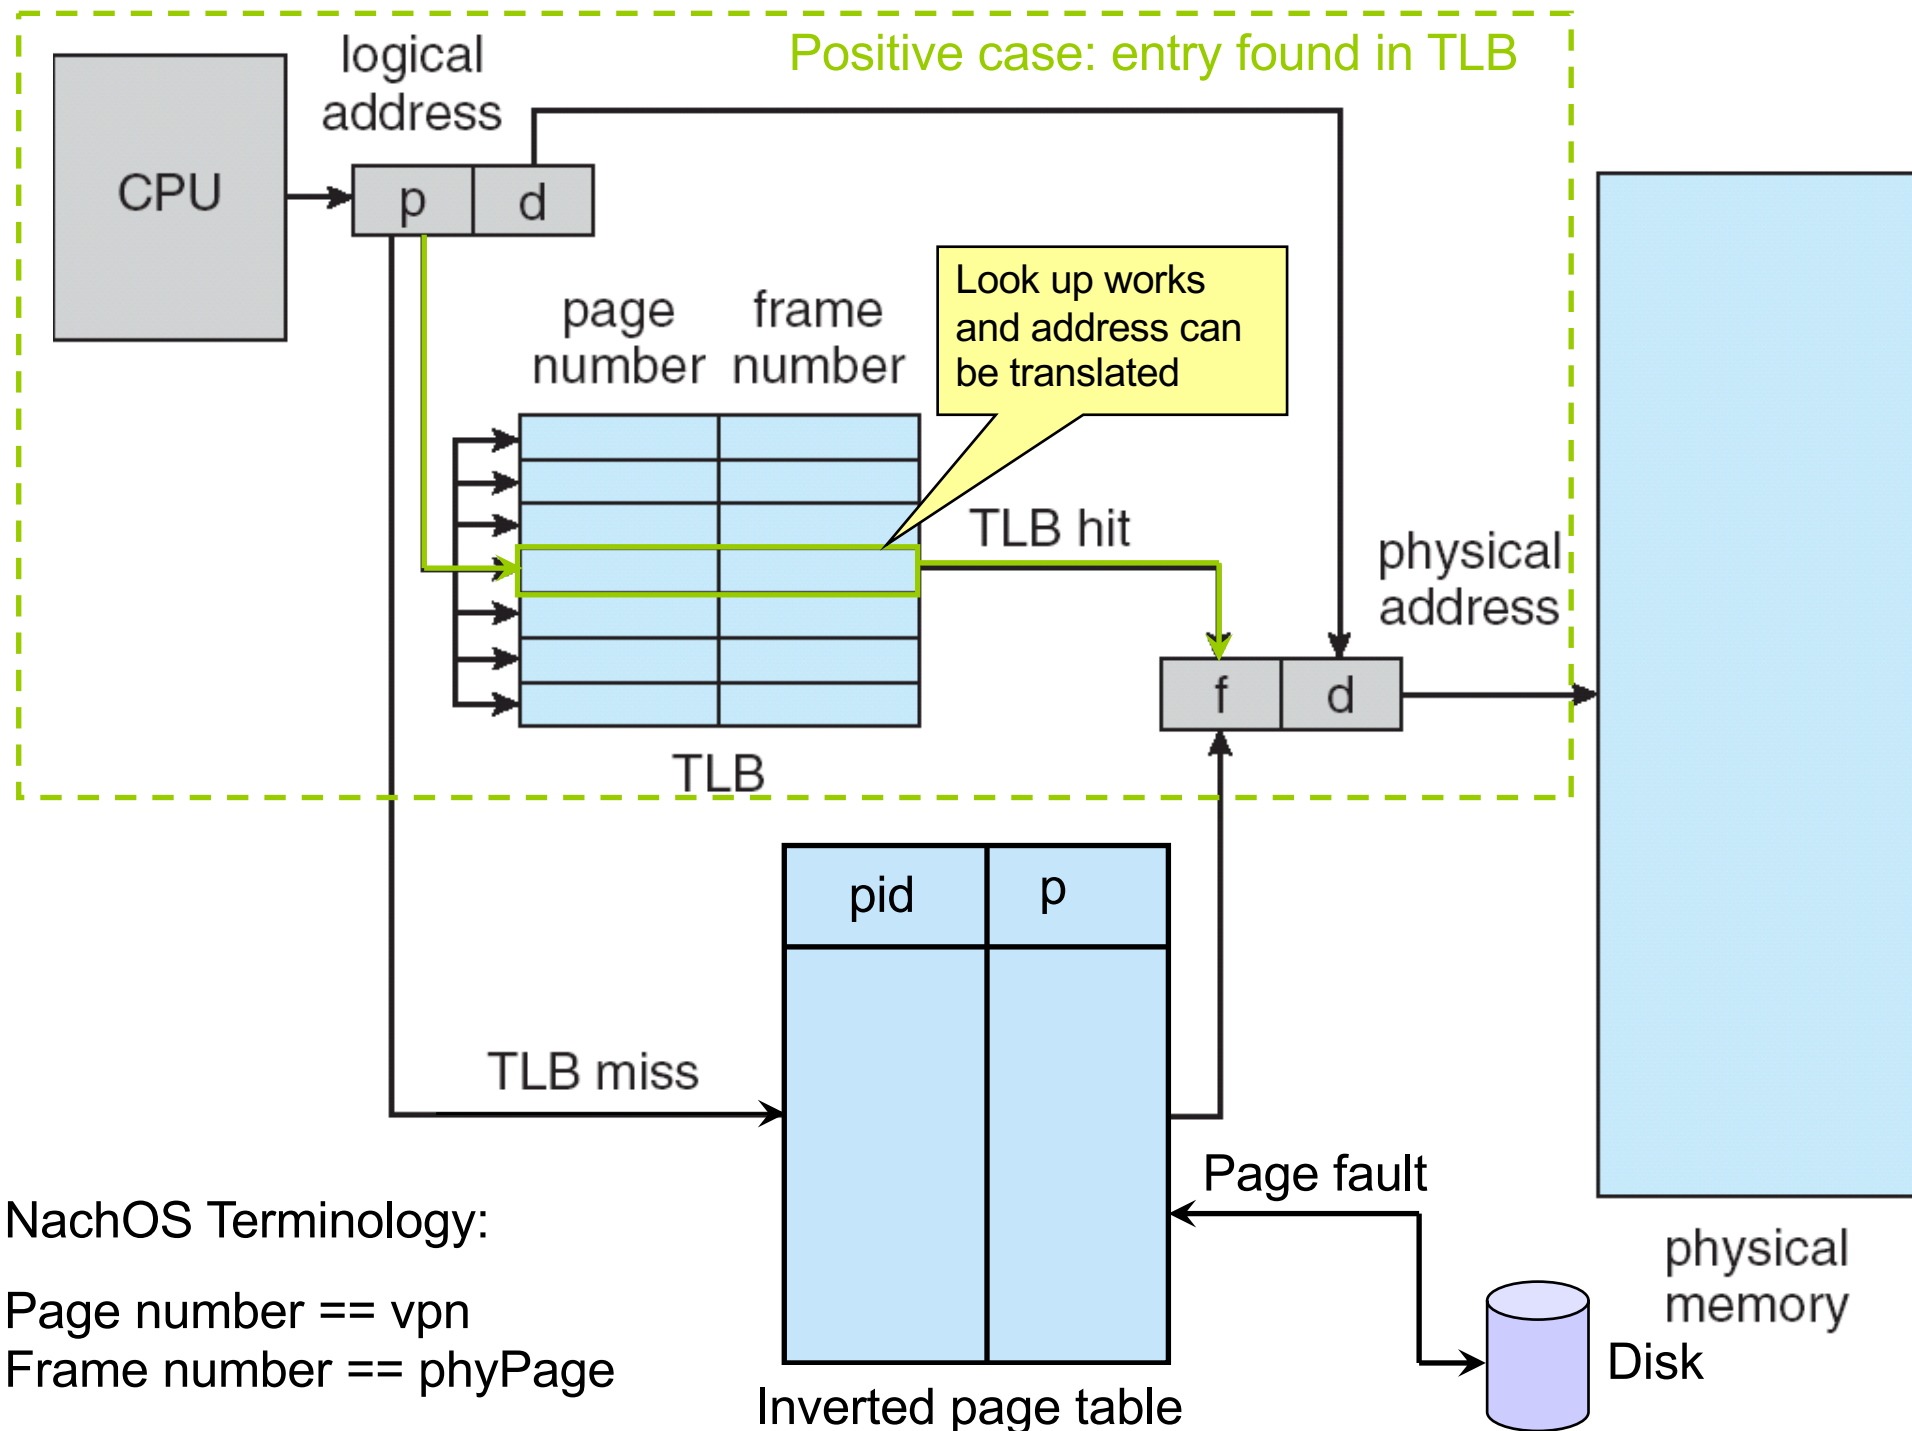




Address Translation in NachOS

Address Translation in NachOS


- How to translate a virtual address into a physical address?
 - Translation Look-aside Buffer (TLB): one per process
 - Inverted Page Table (IPT): one for the entire system
- Function: `Translate(vAddress, ...)`
 - Calculate VPN (Virtual Page Number) and offset
 - Lookup VPN in TLB
 - If lookup was successful:
 - Calculate physical address
 - If lookup was not successful (i.e., TLB miss)
 - Update TLB
 - Try same lookup again (this time it should work)



NachOS Terminology:

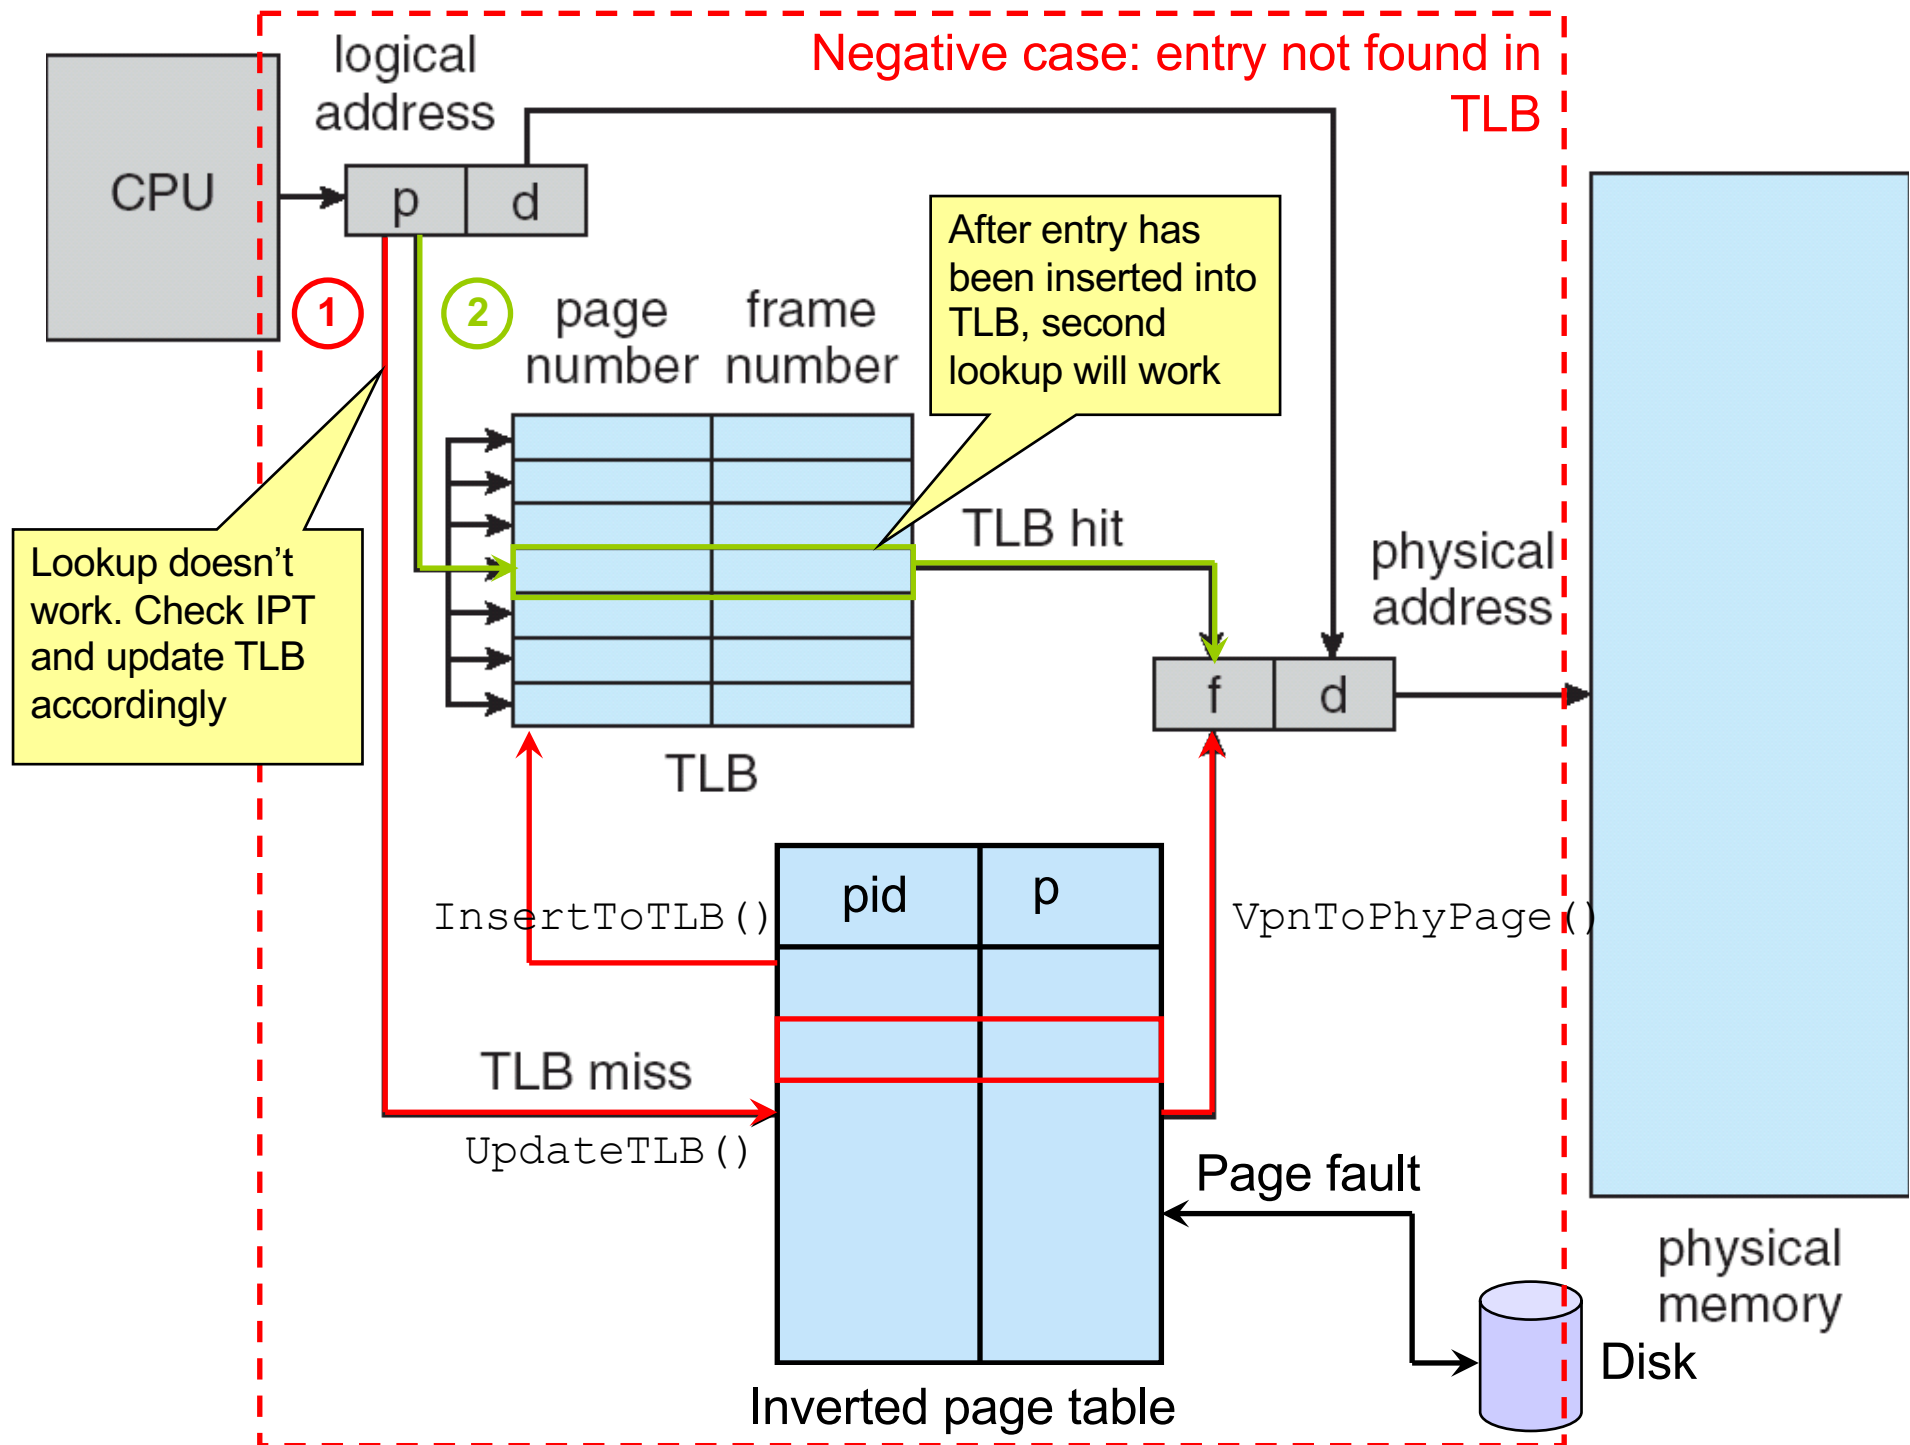
Page number == vpn

Frame number == phyPage



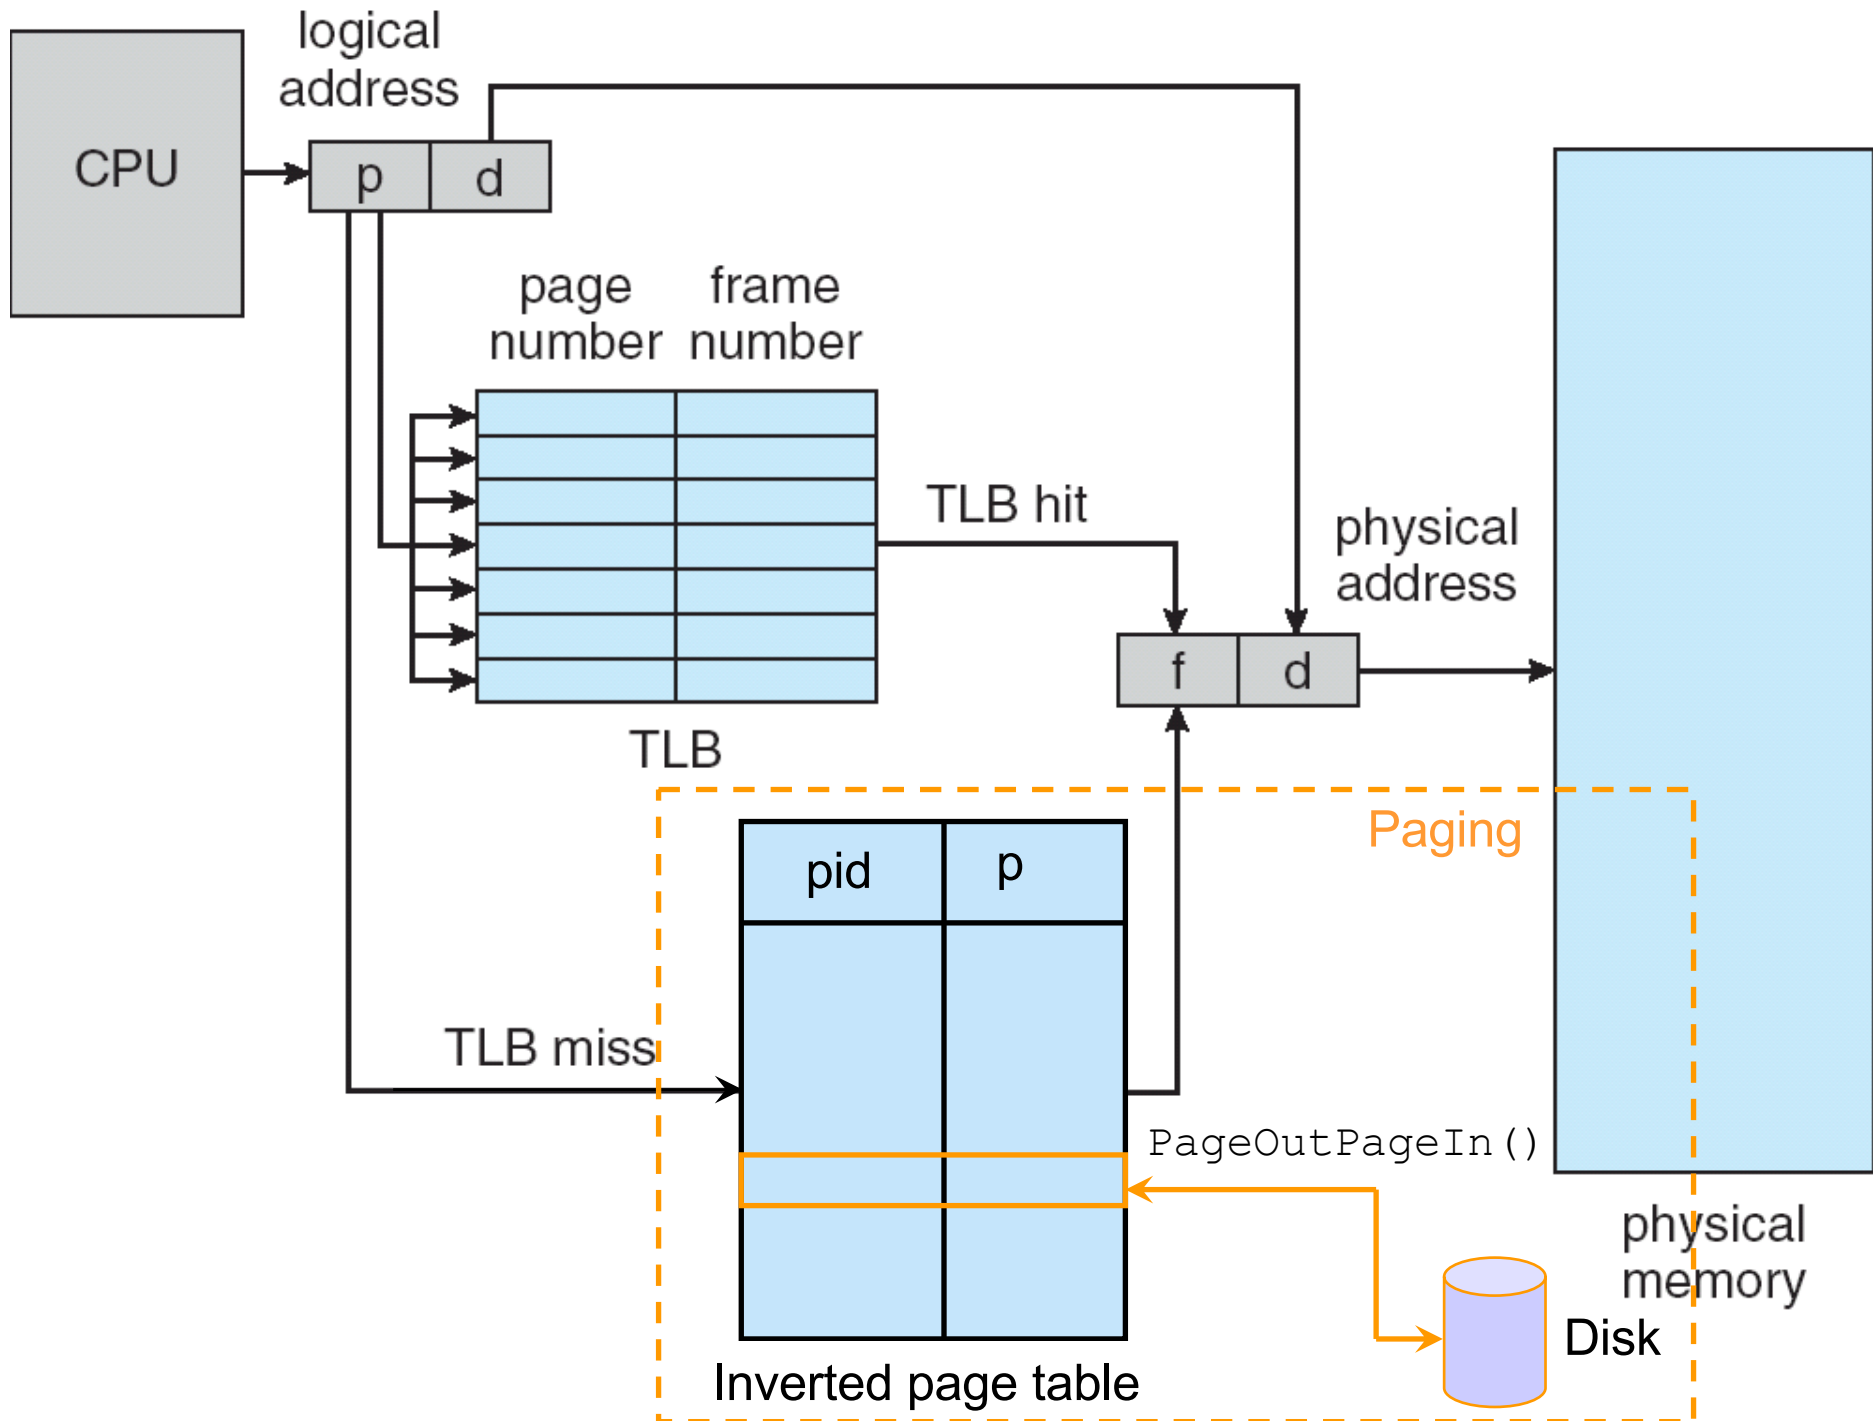
Experiment 3 – Update TLB

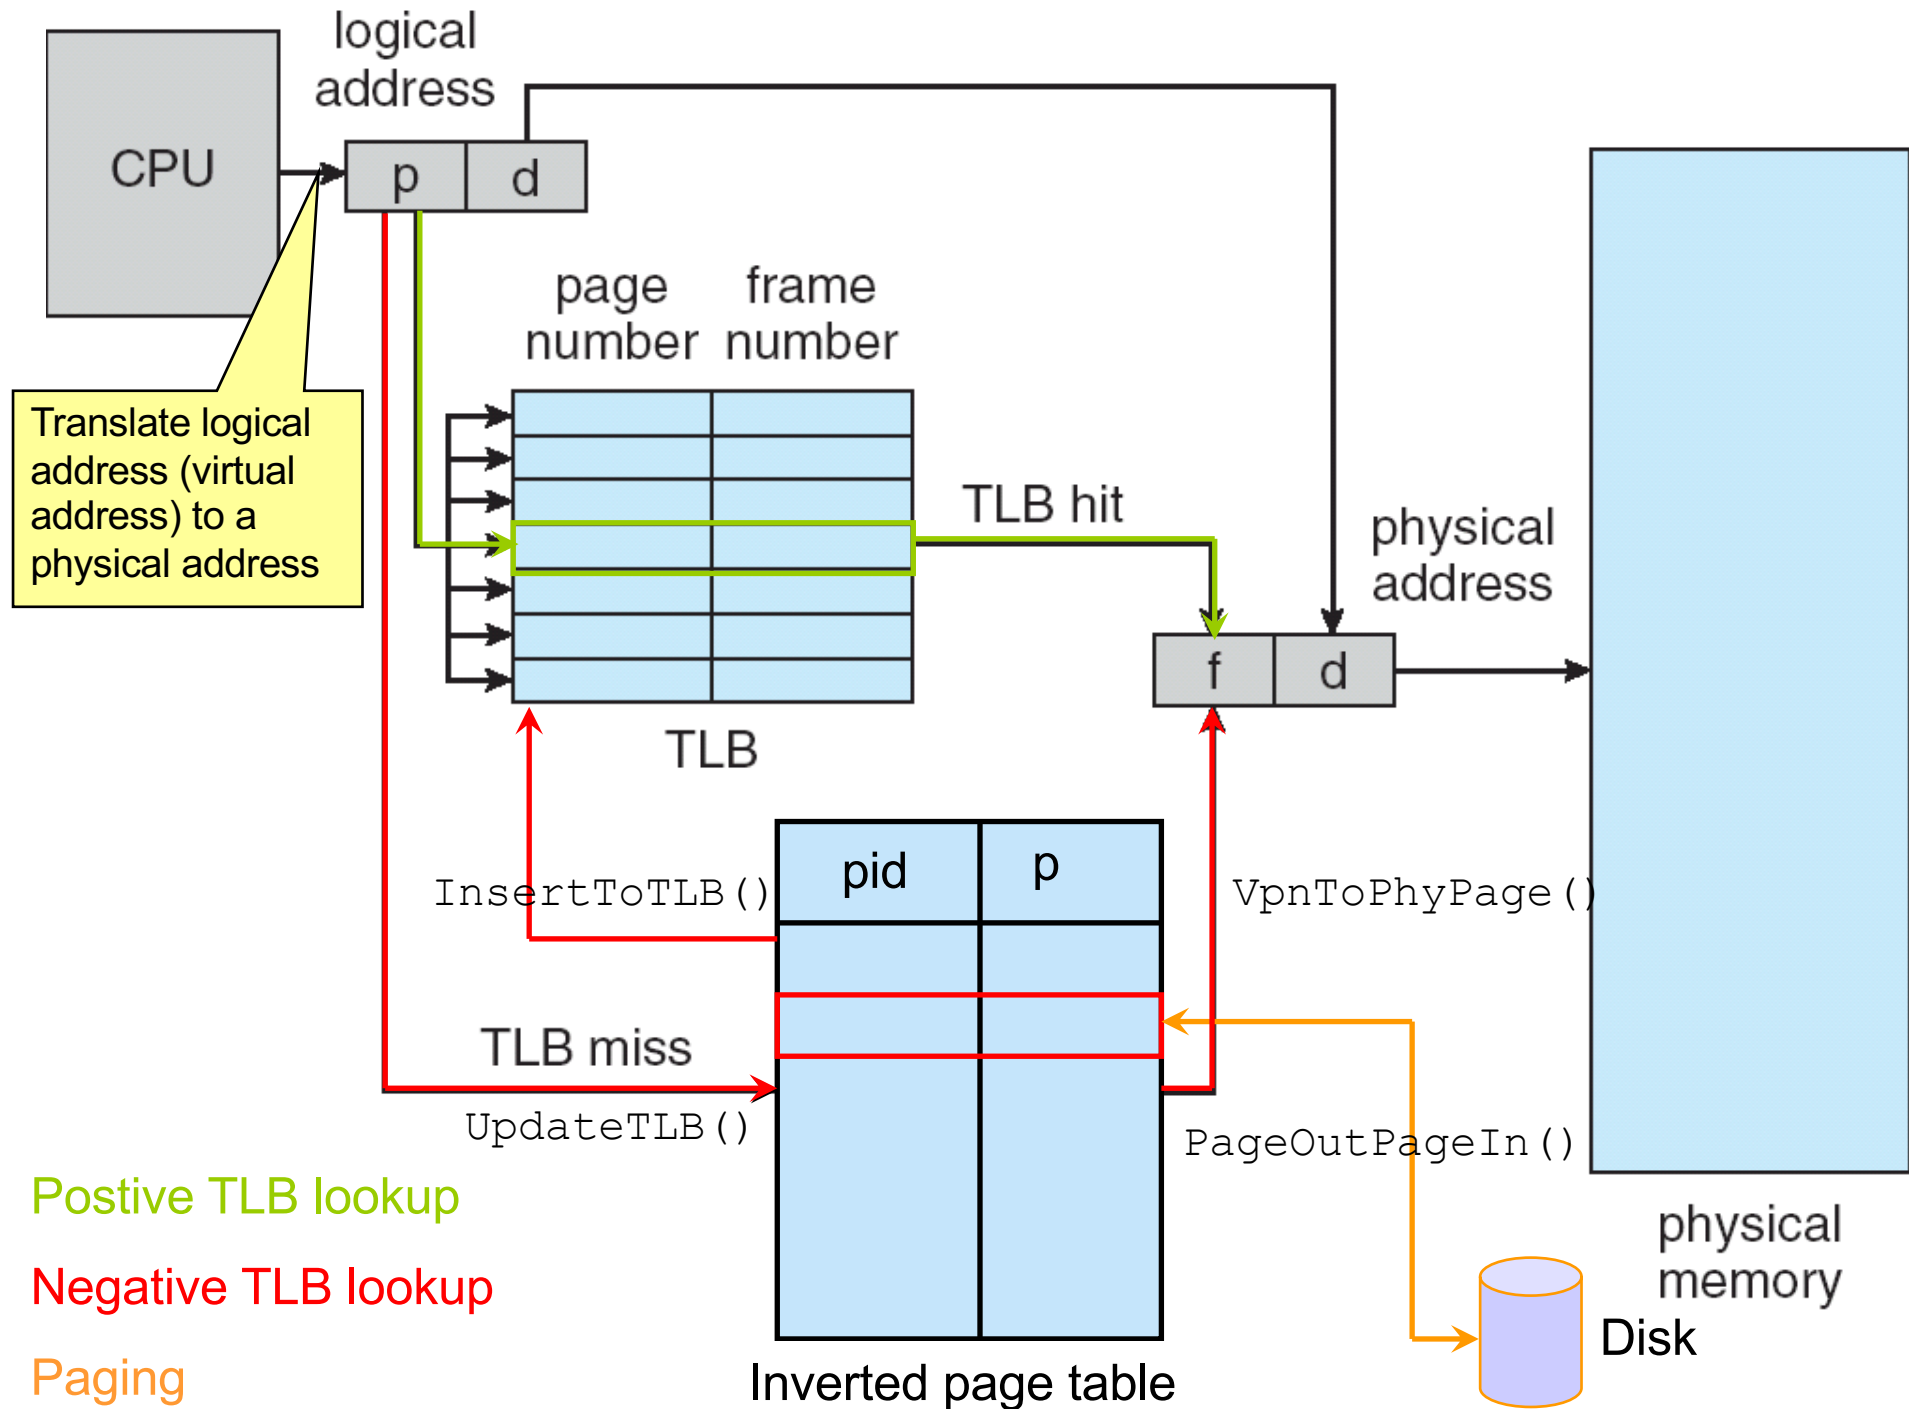
- In case the lookup was not successful (i.e., in case of a TLB miss), how to update the TLB?
- **Function:** `UpdateTLB (vAddress)`
 - Determine VPN based on virtual address
 - Check whether VPN can be found in IPT
 - If VPN can be found:
 - Insert the VPN/PhyPage into TLB
 - If not:
 - Perform paging




Experiment 3 – Update TLB

- If VPN cannot be found in IPT, how to perform paging?
- **Function:** `PageOutPageIn (vpn)`
 - Determine the victim frame using the clock algorithm
 - Page out victim page
 - Write victim page to swap file
 - Page in the new page
 - Load new physical page from swap file
 - Update the IPT table with VPN/PhyPage combination







Discussion of Experiment 3



Experiment 3 – Overview

- Objective

- Understand why TLB can provide fast address translation
- Know how address translation is done by using IPT
- Understand how page replacement is essential to virtual memory
- Understand how to implement a clock replacement algorithm

- Tasks

- Implement missing functionality for address translation
- Run test program and analyse output

Directory Structure

`bin`

For generating NachOS format files, **DO NOT CHANGE!**

`filesystem`

NachOS kernel related to file system, **DO NOT CHANGE!**

`lab1`

Experiment 1, no coding is required.

`lab2`

Experiment 2, process synchronization.

`machine`

MIPS H/W simulation, **DO NOT CHANGE** unless asked.

`Makefile.common`

`Makefile.dep`

For compilation of NachOS,
DO NOT CHANGE!

`network`

NachOS kernel related to network, **DO NOT CHANGE!**

`port`

Additional experiment for students registered with CPE205/CSC205

`readme`

Short description of OS labs and assessments

`test`

NachOS format files for testing virtual memory, **DO NOT CHANGE!**

`threads`

NachOS kernel related to thread management, **DO NOT CHANGE!**

`userprog`

NachOS kernel related to running user applications, **DO NOT CHANGE!**

`vm`

Experiment 3, coding virtual memory (TLB, page replacement)



Experiment 3 – Overview

- Incomplete VM code can be found in
 - `vm/tlb.cc`
- You need to add your code to these 3 functions:
 - `int VpnToPhyPage(int vpn)`
 - `void InsertToTLB(int vpn, int phyPage)`
 - `int clockAlgorithm(void)`
- Binary test program for Experiment 3 is provided
 - Execute the test program:
 - `./nachos -x ../test/vmtest.noff -d`

Experiment 3 – InsertToTLB()

- Put a virtual page number and its associated physical page into the TLB

```
void InsertToTLB(int vpn, int phyPage) {  
    int i = 0; //entry in the TLB  
  
    //your code to find an empty in TLB or to  
    //replace the oldest entry if TLB is full  
    ???  
  
    //copy dirty data to memoryTable  
    ...  
    memoryTable[phyPage].clockCounter = 0  
}
```

This is the reset to 0, mentioned in the lab manual. It's already done for you, so you don't need to worry about it

- What do you need to do?
 - Replace the ??? with your code, don't touch the rest
 - Your code has to make sure that *i* is set correctly
 - i.e., the *i*-th entry in the TLB is either *empty* or the *oldest*

Experiment 3 – InsertToTLB()

- Check all entries in the TLB whether there are any invalid entries
- How to do that?
 - The TLB is an array of `TranslationEntry` objects:
 - `TranslationEntry *tlb;`
 - You can access the TLB in the following way:
 - `machine->tlb[i]`
 - A translation entry has several flags
 - For example: you can check whether an entry is valid:
 - `machine->tlb[i].valid`
 - Check `machine/translate.h` for more details
 - The size of that array is defined by the constant `TLBSize`

Experiment 3 – InsertToTLB()

- If there is an invalid TLB entry, then i should point to it
- For example, if the 2nd entry is invalid then $i=1$
 - The new VPN/PhyPage value will be inserted into the 2nd entry of the table

# Entry	VPN	PhyPage	Valid	ReadOnly	Use	Dirty
0	TRUE
1	FALSE
2	TRUE

Experiment 3 – InsertToTLB()


- If there is no invalid entry, then i should point to the oldest entry
 - The oldest entry will be replaced by the new VPN/PhyPage values
 - You need to keep track of the oldest entry
 - C++ hint: use a static variable for that purpose
 - `static int FIFOPointer = 0;`
 - Once a static variable is initialised, it remains in the memory
 - No re-initialisation afterwards
 - Make sure FIFOPointer is always correctly pointing to the oldest entry
 - Simple FIFO: if an entry is just inserted, then the entry next to it is the oldest entry
 - `FIFOPointer = (i + 1) % TLBSize`

Experiment 3 – VpnToPhyPage()

- Return the physical page for a VPN (if it exists in the IPT)

- ```
int VpnToPhyPage(int vpn) { ... }
```

- IPT is realised as hash table:



**Access to the first entry:**

```
IpEntry* iptPtr = hashIPT(vpn, currentThread->pid)
```

- Multiple entries are chained in a linked list
  - To access the next element use `iptPtr = iptPtr->next`
- Return the `iptPtr->phyPage` value from the IPT entry for which the following condition is true:
  - `iptPtr->vPage==vpn && pid==currentThread->pid`
- Return -1 if no entry can be found that matches the above condition

# Experiment 3 – clockAlgorithm()

- Determines which physical page should be paged out
  - `int clockAlgorithm(void)`
- A memory table is used by the clock algorithm
  - Can be accessed by using `memoryTable[i]`
- This memory table has as many entries as there are physical pages
  - A constant is used for that purpose: `NumPhysPages`
- Three things are of importance
  - `memoryTable[i].valid` – Is the *i*-th entry valid?
  - `memoryTable[i].dirty` – Is the *i*-th entry dirty?
  - `memoryTable[i].counter` – How often has the page been used?

# Experiment 3 – clockAlgorithm()

- A `clockPointer` is used which has to remember its position
  - Use a static variable for this purpose
    - `static int clockPointer = 0;`
- Repeatedly loop over all entries until you find a page that fulfills one of these conditions:
  - Is invalid:
    - `memoryTable[clockPointer].valid == false`
  - Is not dirty and old enough:
    - `!memoryTable[clockPointer].dirty && memoryTable[clockPointer].clockCounter == OLD_ENOUGH`
  - Is dirty, old enough, and already got a second chance
    - `memoryTable[clockPointer].dirty && memoryTable[clockPointer].clockCounter == OLD_ENOUGH + DIRTY_ALLOWANCE`

# Experiment 3 – clockAlgorithm()

- After you have checked these conditions for one entry, increase the clock counter
  - `memoryTable[clockPointer].clockCounter++`
  - This is to ensure that eventually one page will be old enough
- Also, in order to check the next page, advance the clock pointer
  - `clockPointer = (clockPointer + 1) % NumPhysPages`
- Once a victim has been found
  - Return the page number (i.e., return `clockPointer`)
  - Advance the clock pointer before return (important because the page that is currently being pointed at will be replaced)



# Experiment 3 – Analysis of Output

| VPN | TLB Miss/<br>Page Fault | Page In | Page Out | PhyPage | TLBEntry<br>Inserted |
|-----|-------------------------|---------|----------|---------|----------------------|
| 0   | Page Fault              | 0       | -        | 0       | 0                    |
| 9   | Page Fault              | 9       | -        | 1       | 1                    |
| ... | ...                     | ...     | ...      | ...     | ...                  |
| 0   | TLB Miss                | -       | -        | 0       | 1                    |
| ... | ...                     | ...     | ...      | ...     | ...                  |
| 9   | Page Fault              | 9       | 27       | 2       | 2                    |

- Analyse the output of your program and complete the table
- Write down the following
  - Page size (defined in NachOS)
  - Number of physical frames (defined in NachOS)
  - TLB size (defined in NachOS)
  - Number of pages used by the test program
  - Number of page faults that occurred during execution of the test program



# Experiment 3 – Summary

- No group effort
- Leave a working copy of your code in your account
- Report:
  - Include analysis of test program output which verifies that your solution works
  - Clearly explain which TLB entry or physical frame should be used whenever there is a TLB miss or page fault
  - Do not attach any debug output
  - Due ONE WEEK AFTER your lab session

# Acknowledgement

- The slides are revised from the previous versions created by Dr. Heiko Aydt.