

Part 2: Processes and Threads

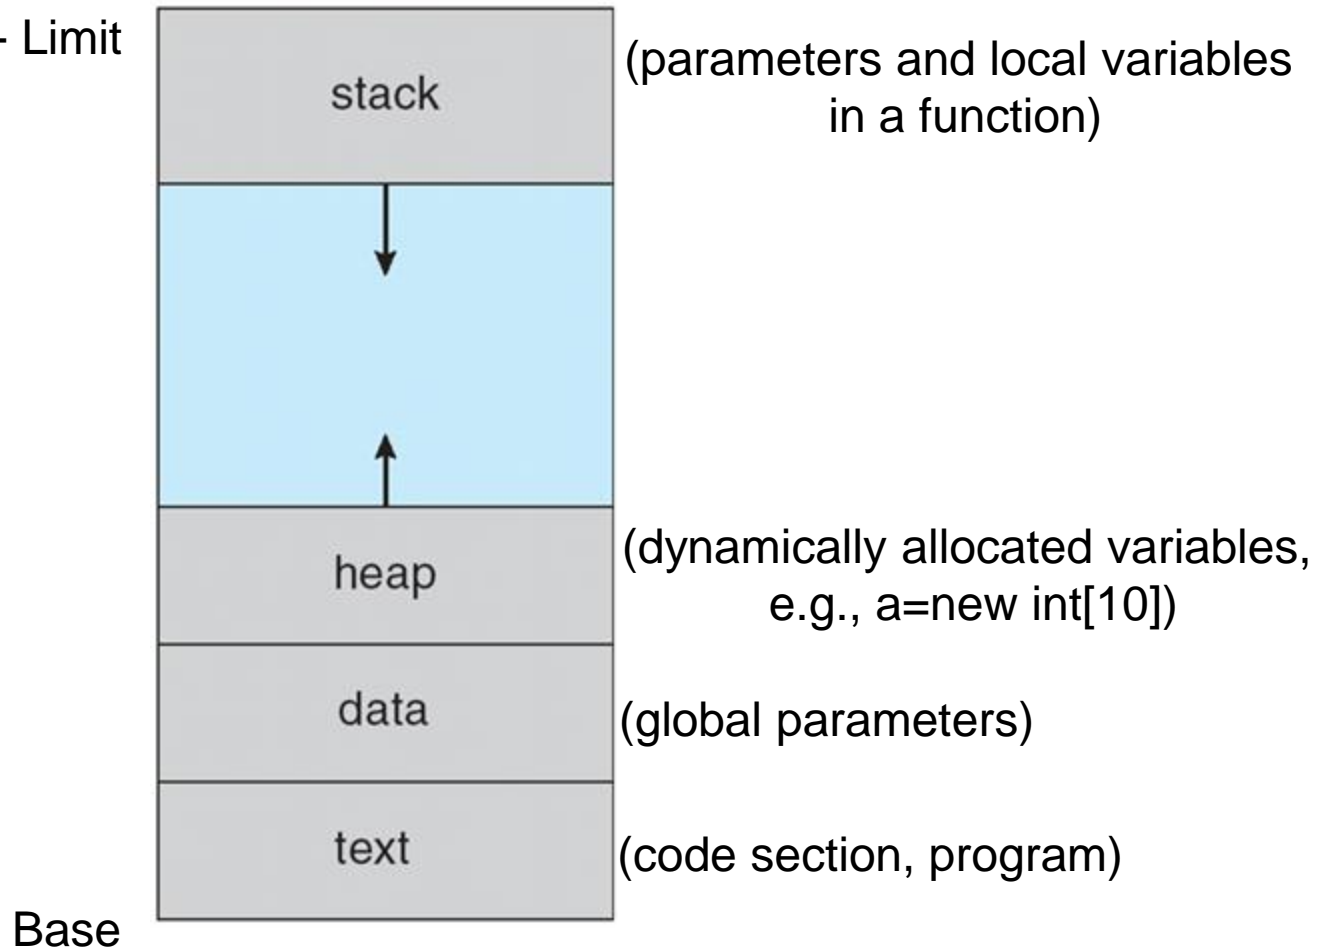
- Process Concept
- Process Scheduling
- Operation on Processes
- Cooperating Processes
- Interprocess Communication
- Threads

Process Concept

- An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-sharing system – user programs & commands
- Textbook uses the terms *job* and *process* almost interchangeably (**Job = Process**).
- Process – a program in execution; process execution must progress in sequential fashion.

Process in Memory

Base + Limit

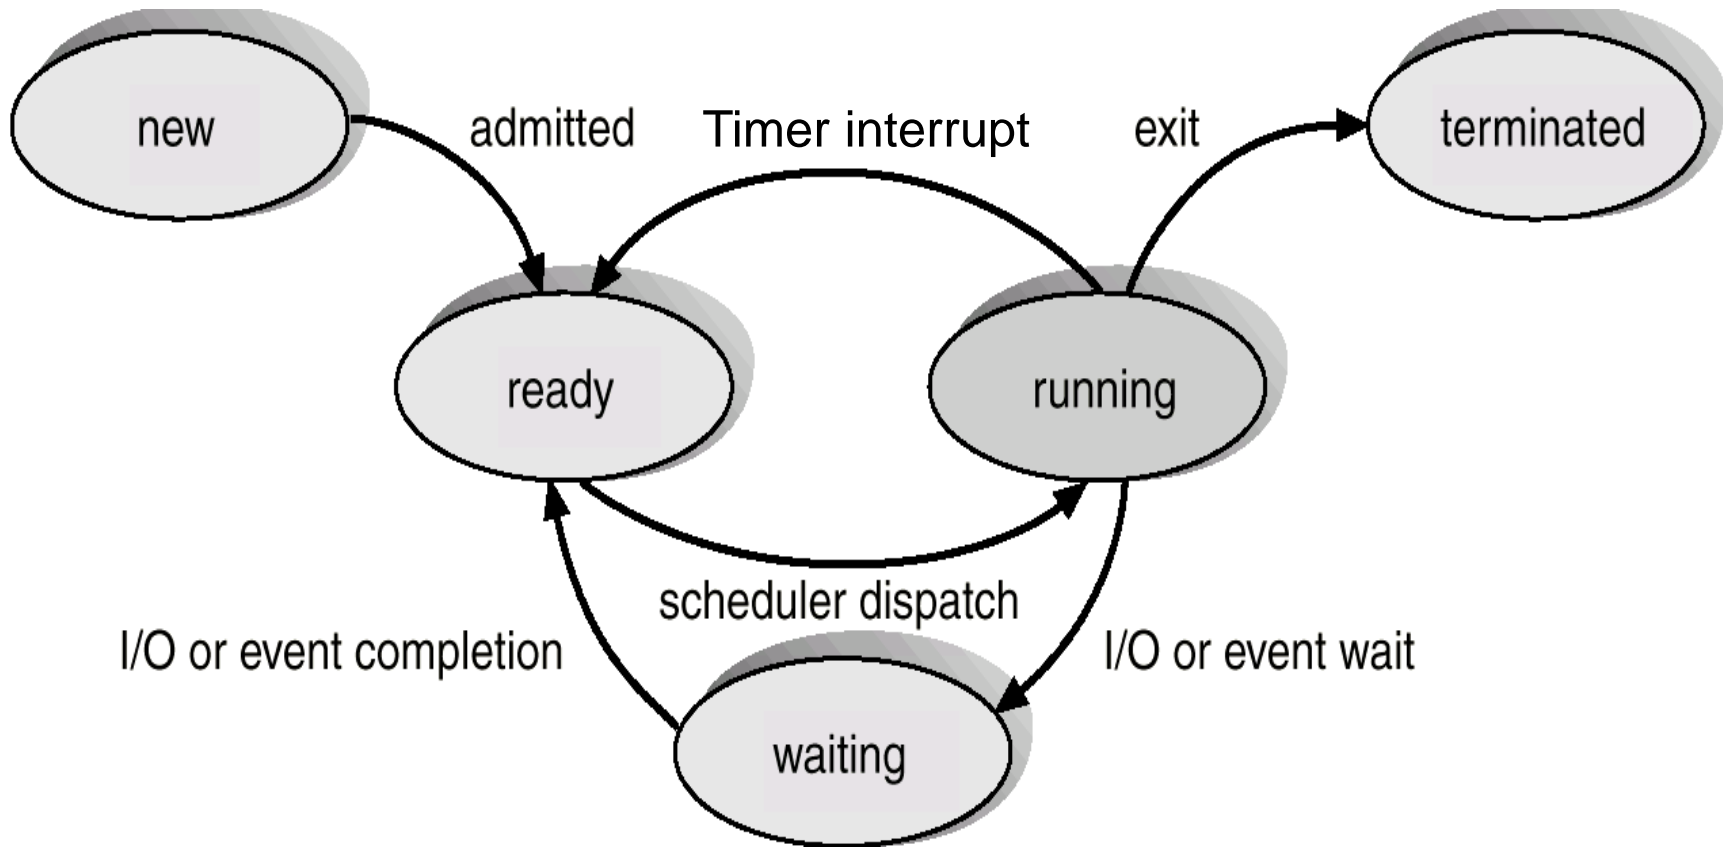


Base

Process State

- As a process executes, it changes *state*
 - new: The process is being created.
 - running: Instructions are being executed.
 - waiting: The process is waiting for some event to occur.
 - ready: The process is waiting to be assigned to CPU.
 - terminated: The process has finished execution.

Diagram of Process State



Process Control Block (PCB)

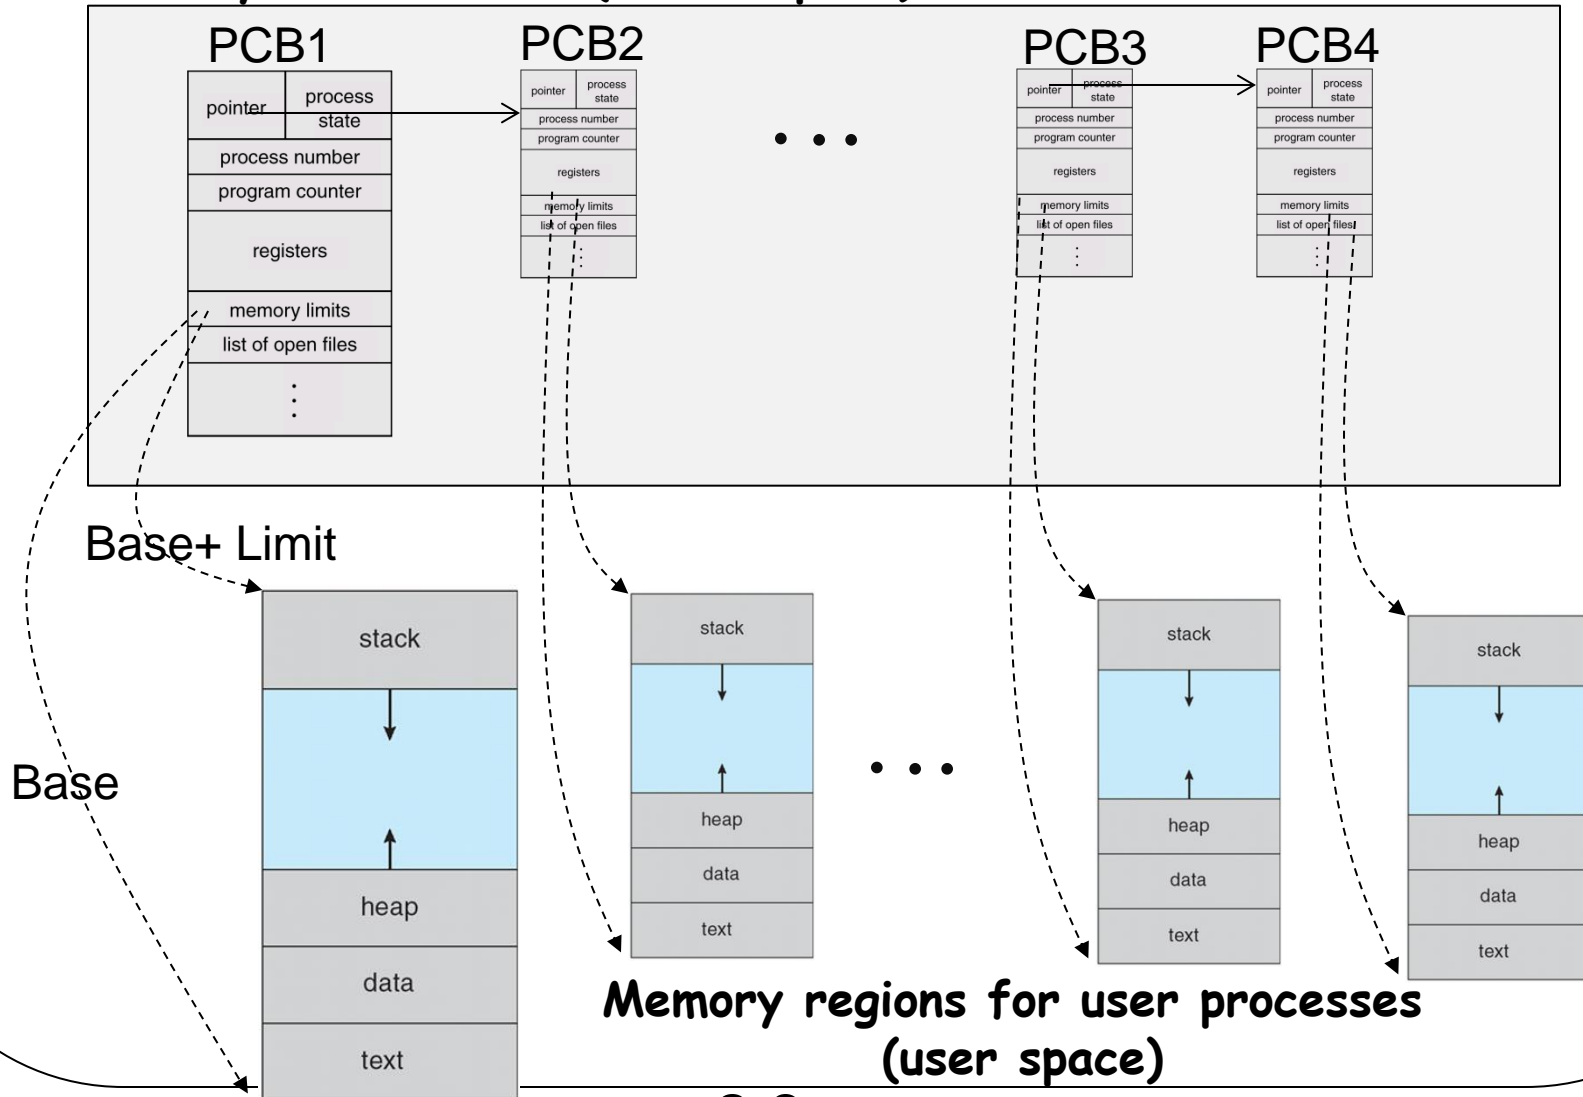
- Keep the information associated with each process:
 - Pointer
 - Process state
 - Process number (process id)
 - Program counter (pointer to the next instruction)
 - CPU registers
 - Memory management information (e.g., memory limit)
 - Information regarding open files (e.g., list of open files)
- In multi-programmed systems, each process has its PCB stored in the main memory.

Process Control Block (PCB)

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

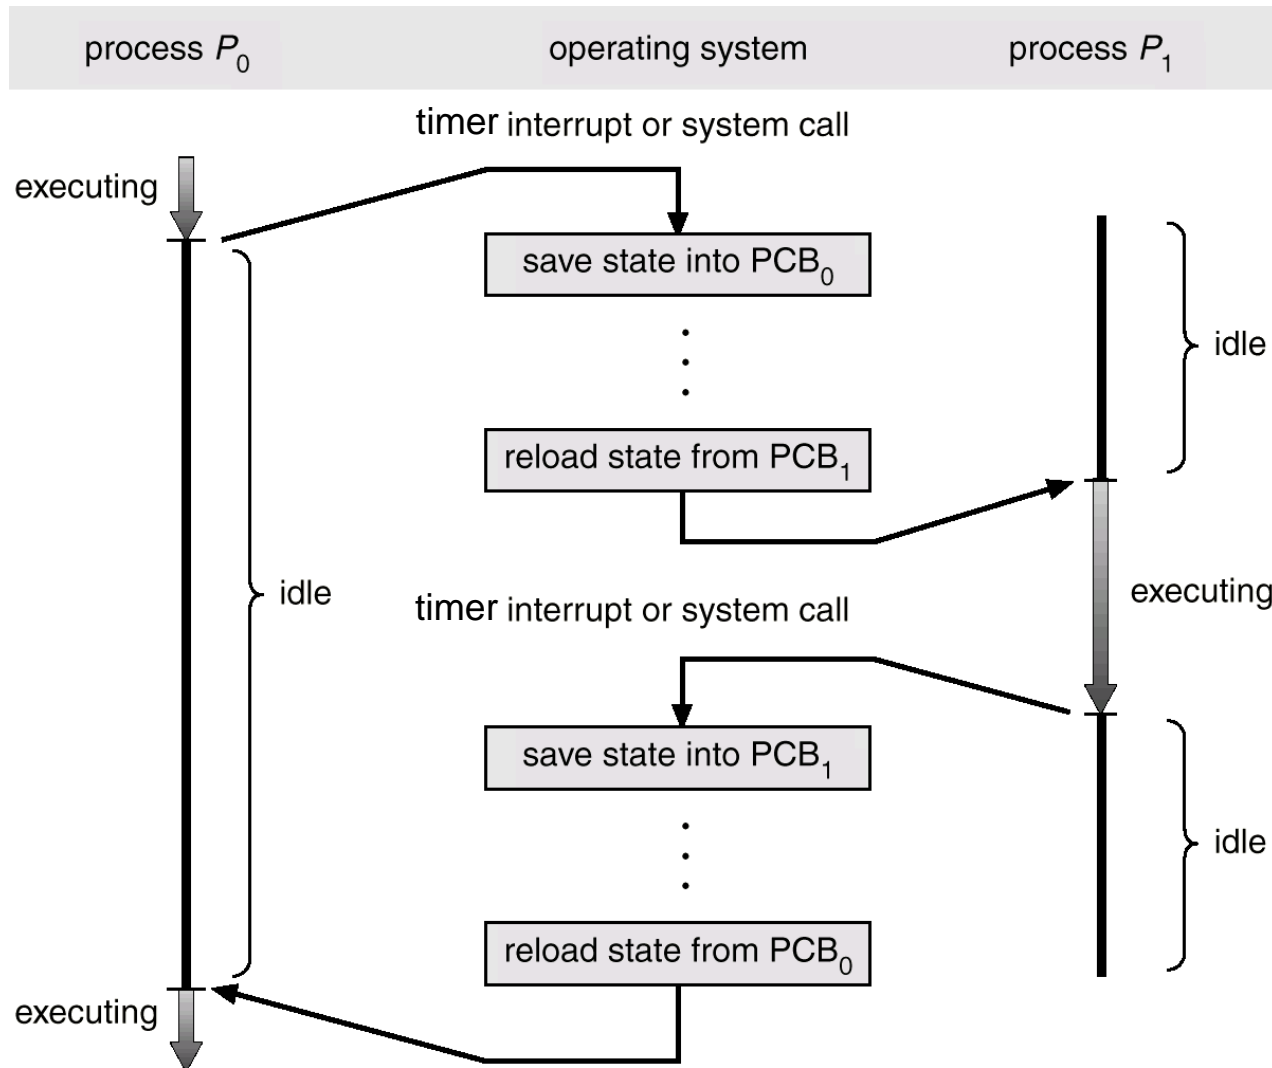
PCB and Processes in Main Memory

Memory area for OS (kernel space)



Memory regions for user processes
(user space)

CPU Switch From Process to Process



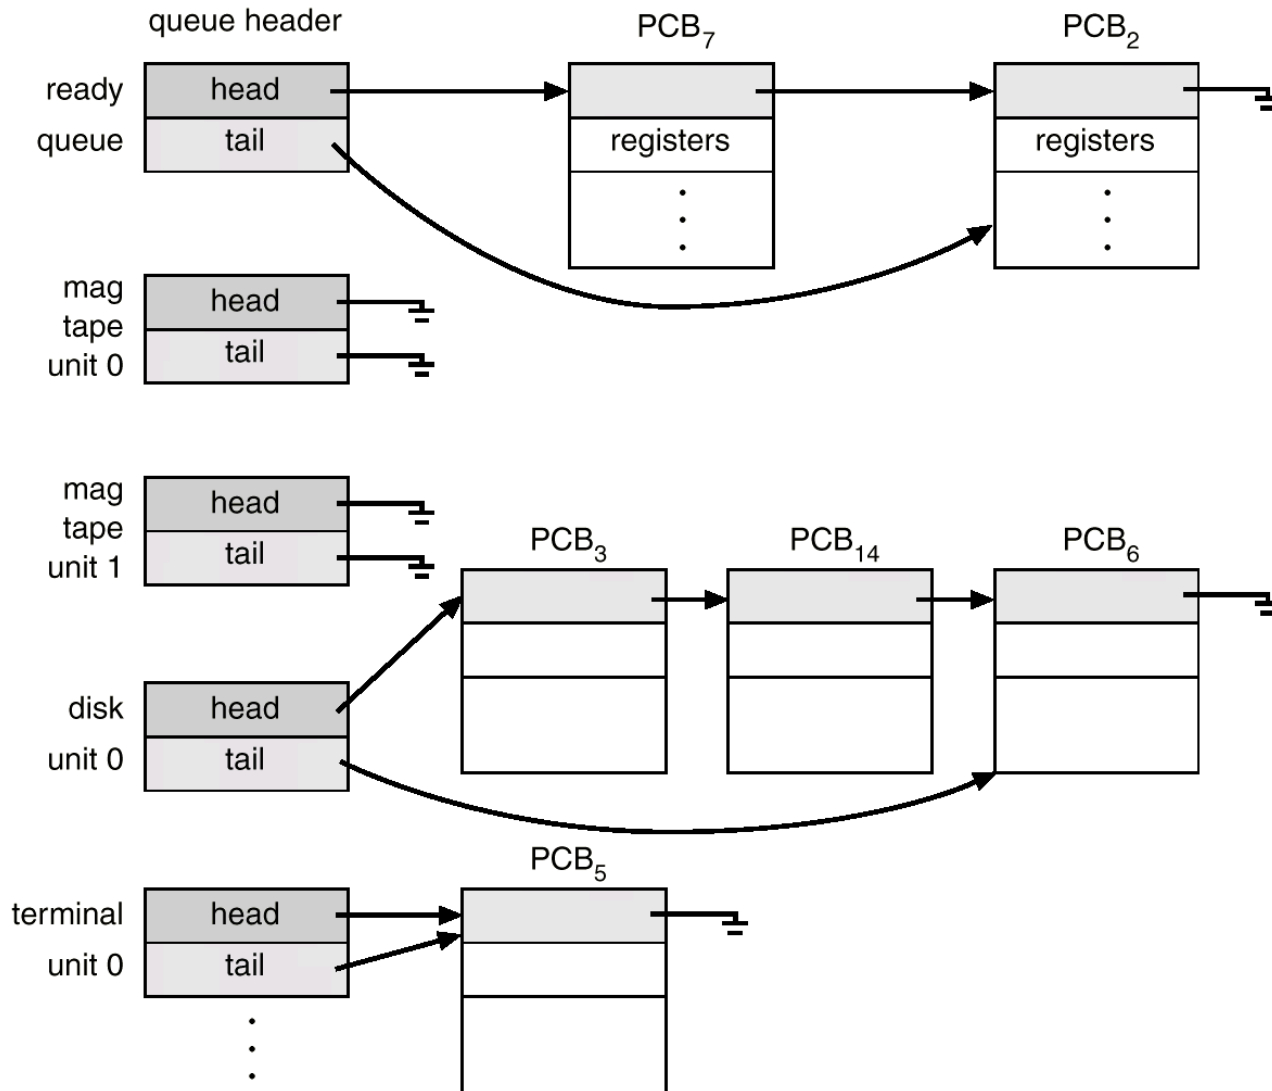
Context Switch

- When CPU switches to another process, the system must save the context (i.e. information) of the old process and load the saved context for the new process.
- Context-switch time is overhead; the system does no useful work while switching.

Process Scheduling Queues

- Job queue – set of all processes *with the same state* in the system.
- Ready queue – set of all processes residing in main memory, ready and *waiting the CPU* to execute.
- Device queues – set of processes *waiting* for an I/O device.
- Process migration between the various queues when process state changes.
- All the queues are stored in the main memory (*kernel space*).

Ready Queue And Various I/O Device Queues



Schedulers

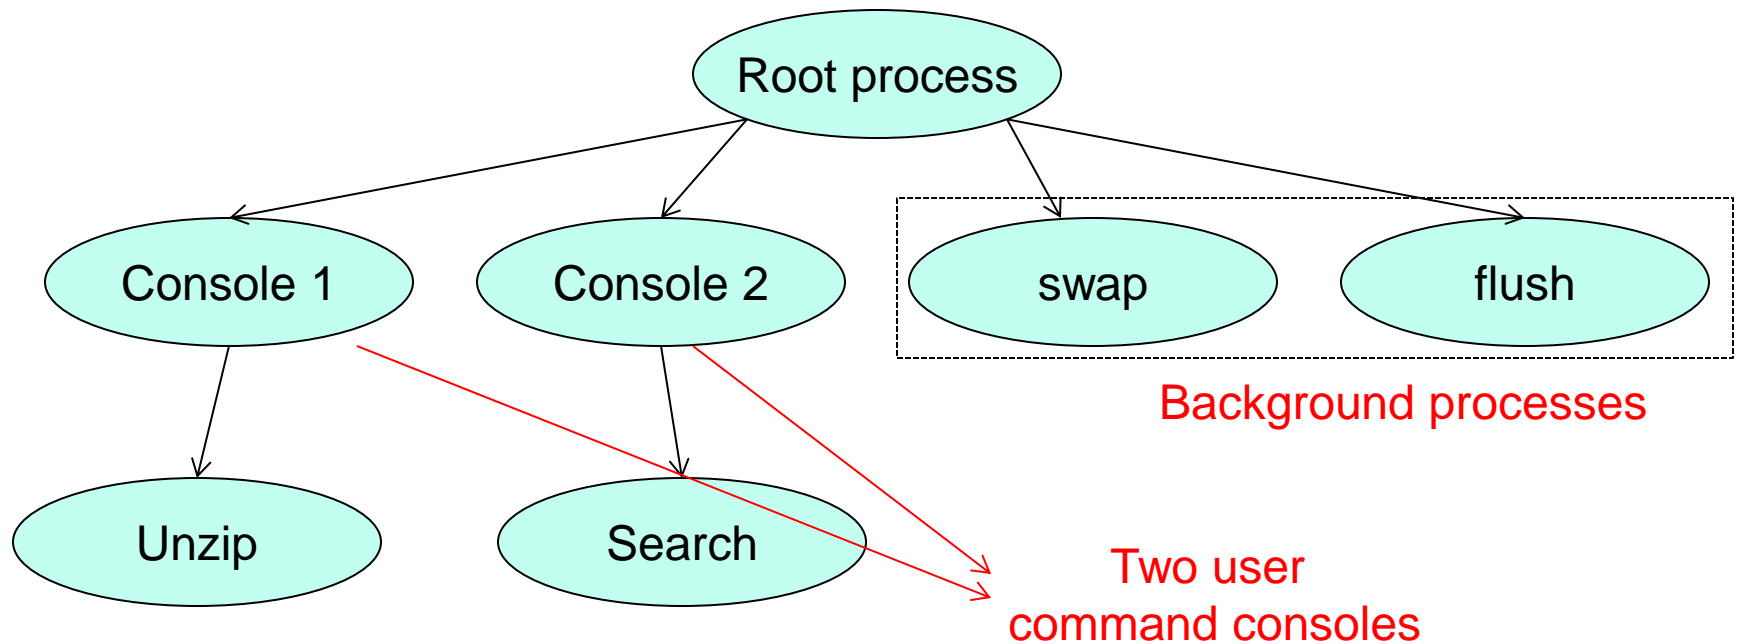
- We need multiple schedulers for different purposes.
- Long-term scheduler (or job scheduler) – selects processes from disk and loads them into memory for execution.
- Short-term scheduler (or CPU scheduler) – selects from among the processes that are ready to execute, and allocate the CPU to one of them.
- Medium-term scheduler
 - When the system load is heavy, **swap** out a partially executed process from memory to hard disk.
 - When the system load is lighter, such processes can be swapped in.

Schedulers (Cont.)

- Short-term scheduler is invoked at least once in every 100 milliseconds because in average a process executes for 100 milliseconds before I/O. It must be fast in making decision.
- Long-term scheduler is invoked very infrequently (seconds, minutes). May be slow in making decision.
- The long-term scheduler controls the *degree of multiprogramming* initially. Then medium-term scheduler after that.

Process Creation

- Parent process creates children processes, which, in turn create other processes, forming a tree of processes (**fork**).



Process Creation (Cont')

- Execution
 - Parent and children execute concurrently.
 - Parent waits until children terminate (**wait**).
- Examples
 - Many Web browsers nowadays fork a new process when we access a new page in a “tab”.
 - OS may create background processes for monitoring and maintenance tasks.


Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**). At that point:
 - A child may output return data to its parent.
 - Process' resources are deallocated by operating system.
- Parent may terminate execution of children processes (**abort**).
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - Parent is exiting.

Cooperating Processes


- *Independent* process cannot affect or be affected by the execution of another process.
- *Cooperating* process can affect or be affected by the execution of another process
- Two models of Interprocess Communication (IPC)
 - Message passing
 - Shared memory

Communication Models



(a)

Message passing



(b)

Shared memory

IPC – Message Passing

- **Message Passing** – processes communicate and synchronize their actions without resorting to shared variables.
- IPC facility provides two operations:
 - **send**(*message*) – message size fixed or variable
 - **receive**(*message*)
- If *P* and *Q* wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive

Direct vs. Indirect Communication

- **Direct communication:** Processes must name each other explicitly:
 - **send** (P , $message$) – send a message to process P
 - **receive**(Q , $message$) – receive a message from process Q
- **Indirect communication:** Messages are sent to or received from mailboxes (also referred to as ports).
 - Mailbox is an object into which messages are placed and removed (like a queue).
 - Each mailbox has a unique id.
 - Primitives are:
 - ***send** (A , $message$): send a message to mailbox A
 - ***receive** (A , $message$): receive a message from mailbox A

Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process.
 - *unbounded-buffer* places no practical limit on the size of the buffer.
 - *bounded-buffer* assumes that there is a fixed buffer size.

Example: Producer & Consumer

//declare a *mailbox* with capacity of B messages

void *producer*(void)

```
{  
    message m;  
    while (1) {  
        //pre-processing...  
        while(mailbox is full) wait;  
        send(mailbox, m);  
    }  
}
```

?

Does this example use direct or indirect communication?

void *consumer*(void)

```
{  
    message m;  
    while (1) {  
        while(mailbox is empty) wait;  
        receive(mailbox, m);  
        //post-processing...  
    }  
}
```

Threads

- This part is for **self-learning**.
- A lecture video will be uploaded to NTU Learn.
- This slide and the later slides in this chapter are **not** examinable.
- About Nachos Labs
 - Nachos uses the term “thread”.
 - All concepts and mechanisms that we learnt about processes can be applicable to Nachos threads.
 - ❄ Thread control block vs. Process control block.
 - ❄ Thread state vs. Process state.
 - ❄ System calls: fork, exit etc.

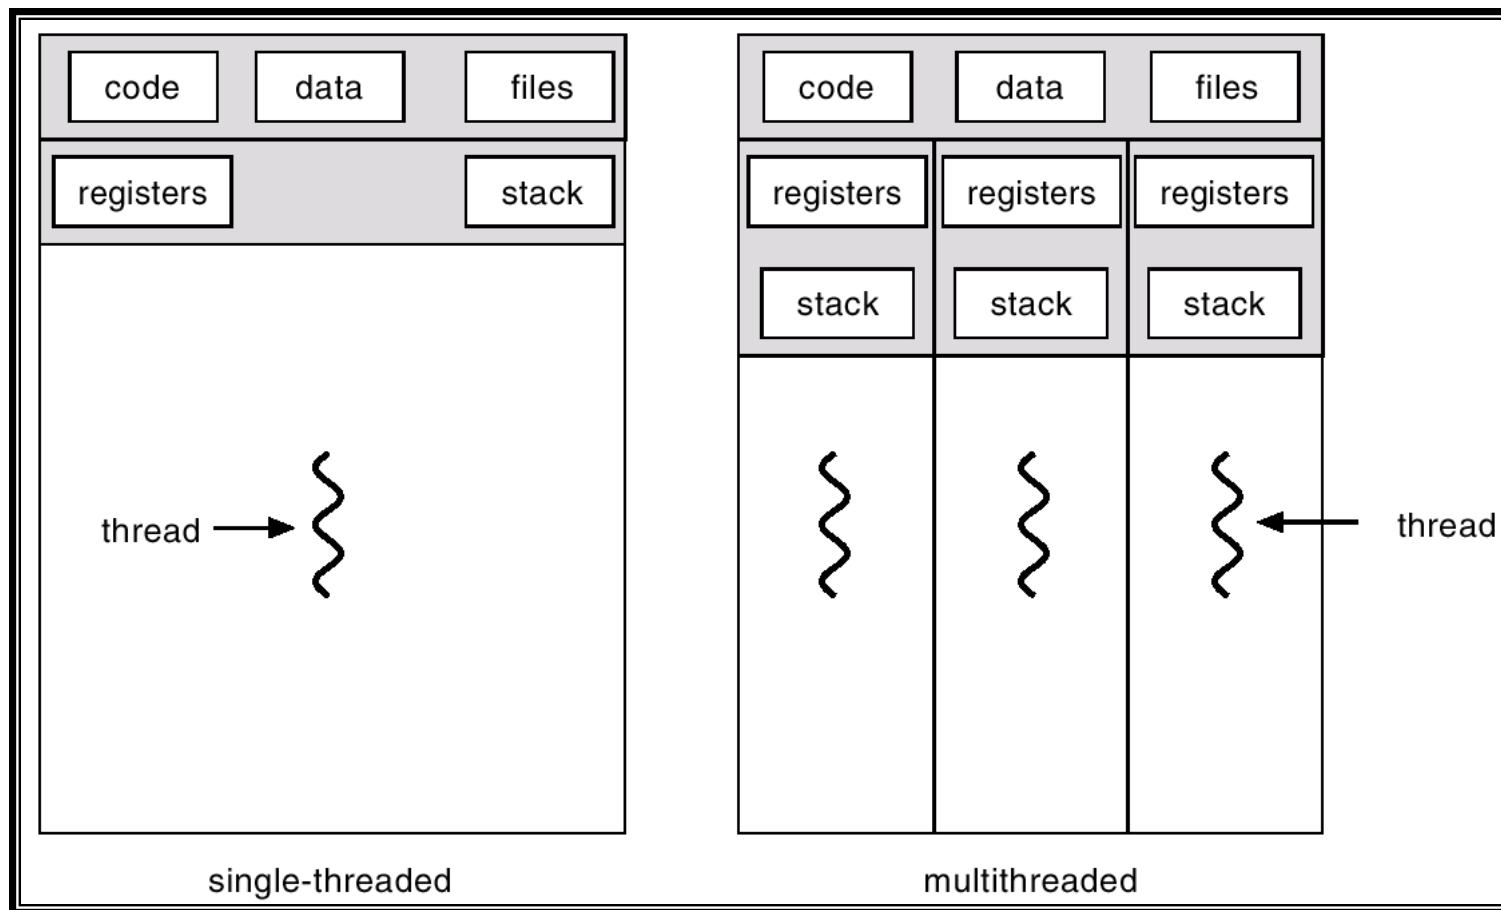
Threads

- Overview
- Single vs Multithreading Process
- Benefits of threads
- Types of threads
- Multithreading models

Overview

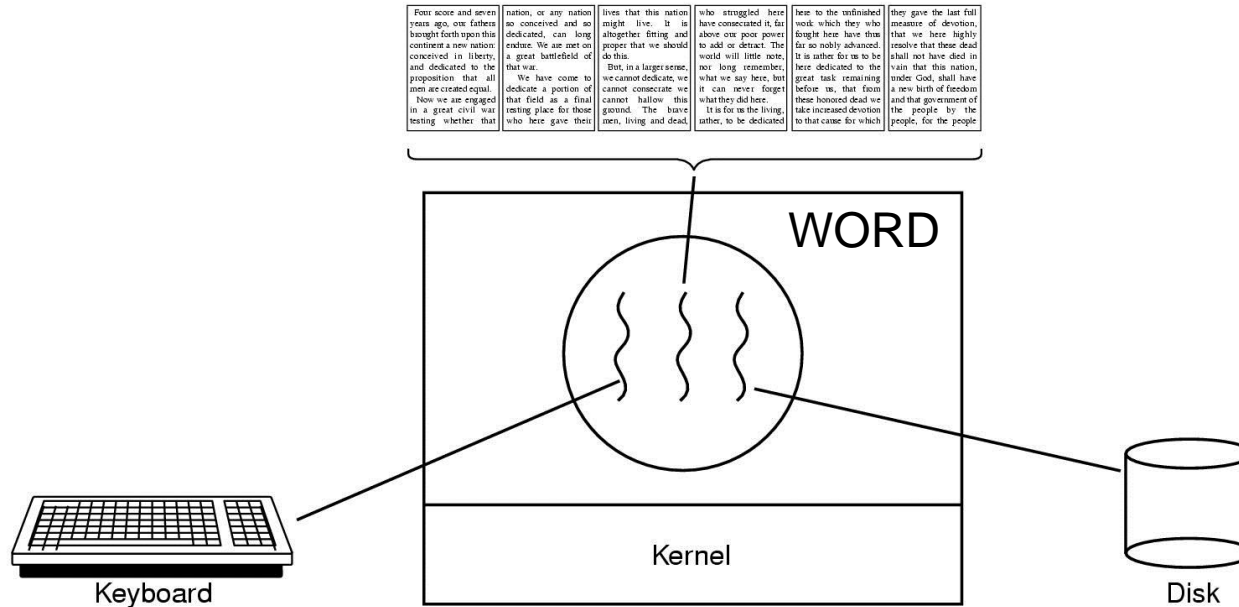
- A *thread* (or *lightweight process*) is a basic unit of CPU utilization; it consists of:
 - a thread id
 - program counter
 - register set
 - stack space
- A thread shares with its peer threads **in the same process** its:
 - code section
 - data section
 - operating-system resources
- A traditional or *heavyweight* process is an executing program with a single thread of control

Single vs Multithreaded Processes

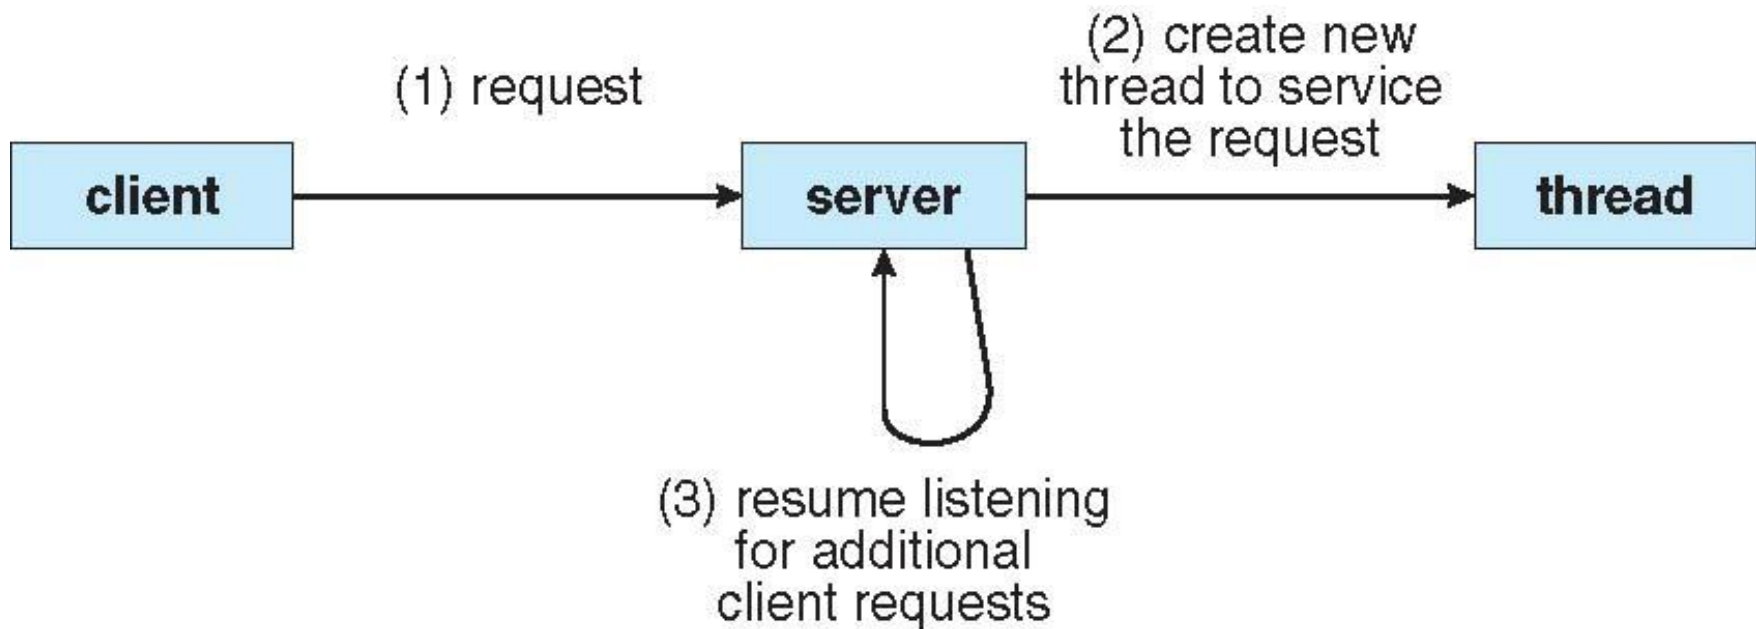


Advantage of Multi-threaded Process

- In a multi-threaded process, while one thread is blocked and waiting, a second thread can run.
 - Cooperation of multiple threads in same process confers higher throughput and improved performance.



Example: Multithreaded Server



Two types of threads

- Paradox
 - Allow users to implement an arbitrary number of threads, **V.S.**,
 - OS kernel can support a limited number of threads due to resource constraints.
- Solution
 - Two layers of abstraction:
 - ✧ User threads (**logical threads**): Supported at the user level; Allow users to create as many threads as they want.
 - ✧ Kernel threads (**physical threads**): Supported at the Kernel level; Slower to create and manage than that of user thread; Resources are eventually allocated in kernel threads.
 - OS maintains the mapping from user threads to kernel threads (**Multithreading models**).

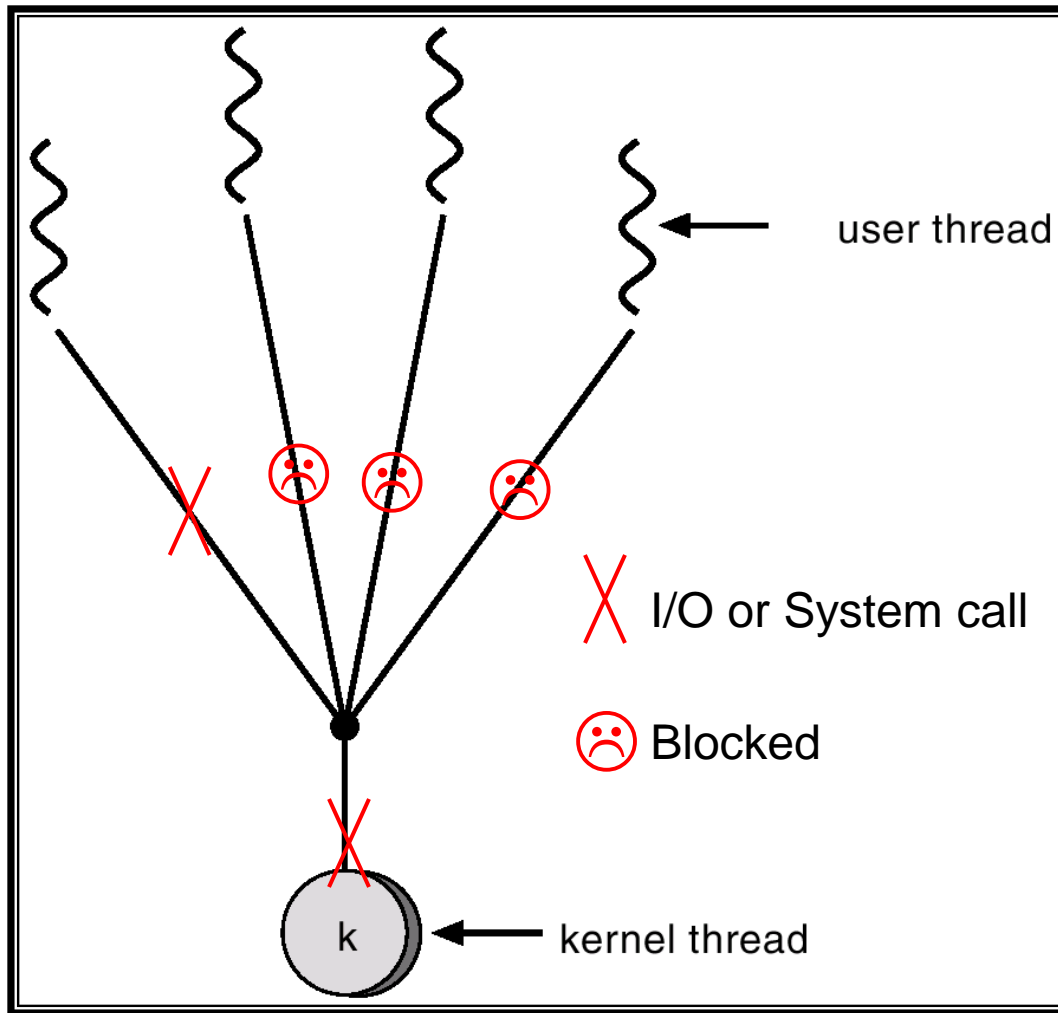
Multithreading Models

- Many-to-One
- One-to-One
- Many-to-Many

Many-to-One

- Many user-level threads mapped to single kernel thread.
- Disadvantage: Unable to run on multiprocessor due to single kernel thread (For those threads mapped to the same kernel thread)


Many-to-One



One-to-One

- Each user-level thread maps to kernel thread.
- Examples: Windows 95/98/NT/2000, OS/2
- Provide more concurrency than the many-to-one model
- Disadvantage:
 - Creating a user thread requires creating the corresponding kernel thread
 - May create too many kernel threads which is a burden to the system (if not so many processors in the system)


One-to-one



Many-to-Many

- Allows many user level threads to be mapped to many kernel threads.
- Allows the operating system to create a sufficient number of kernel threads.
- Examples: Solaris 2, Windows NT/2000
- Do not have the disadvantages of one-to-one and many-to-one model.

Many-to-Many



Advanced Readings

- Beej's Guide to Unix Interprocess Communication, (pdf in EdveNTUre)
 - <http://beej.us/guide/bgipc/>
 - Very good introduction on the IPC implementation inside Unix.
- Tutorials on how to use thread libraries:
 - POSIX pthreads, <http://randu.org/tutorials/threads/>
 - Windows Process and Thread Functions.
[http://msdn.microsoft.com/en-us/library/windows/desktop/ms684847\(v=vs.85\).aspx](http://msdn.microsoft.com/en-us/library/windows/desktop/ms684847(v=vs.85).aspx)