

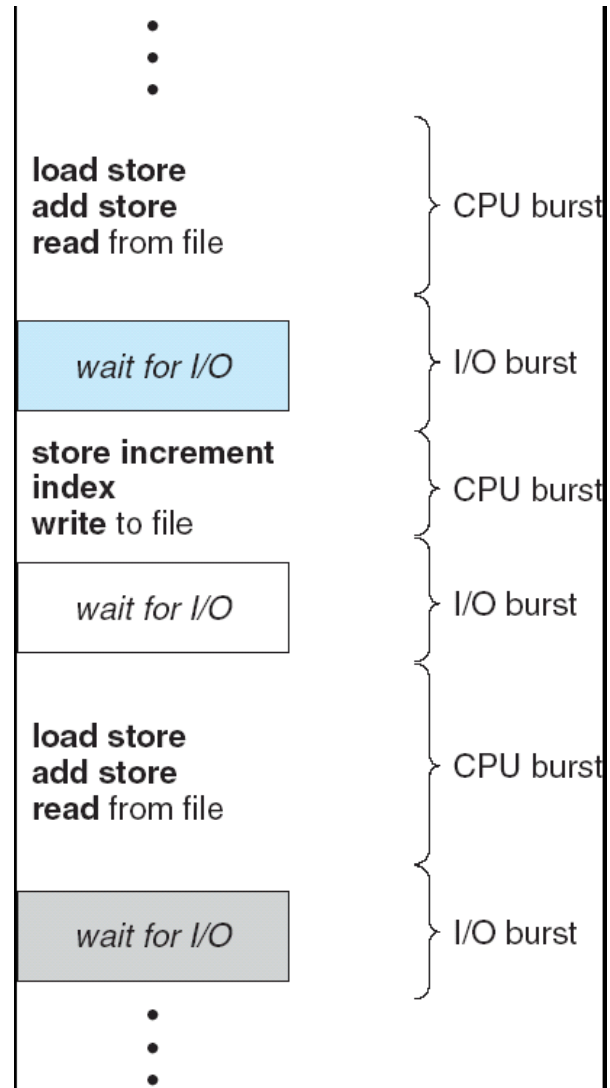
Part 3: Process Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms

Basic Concepts

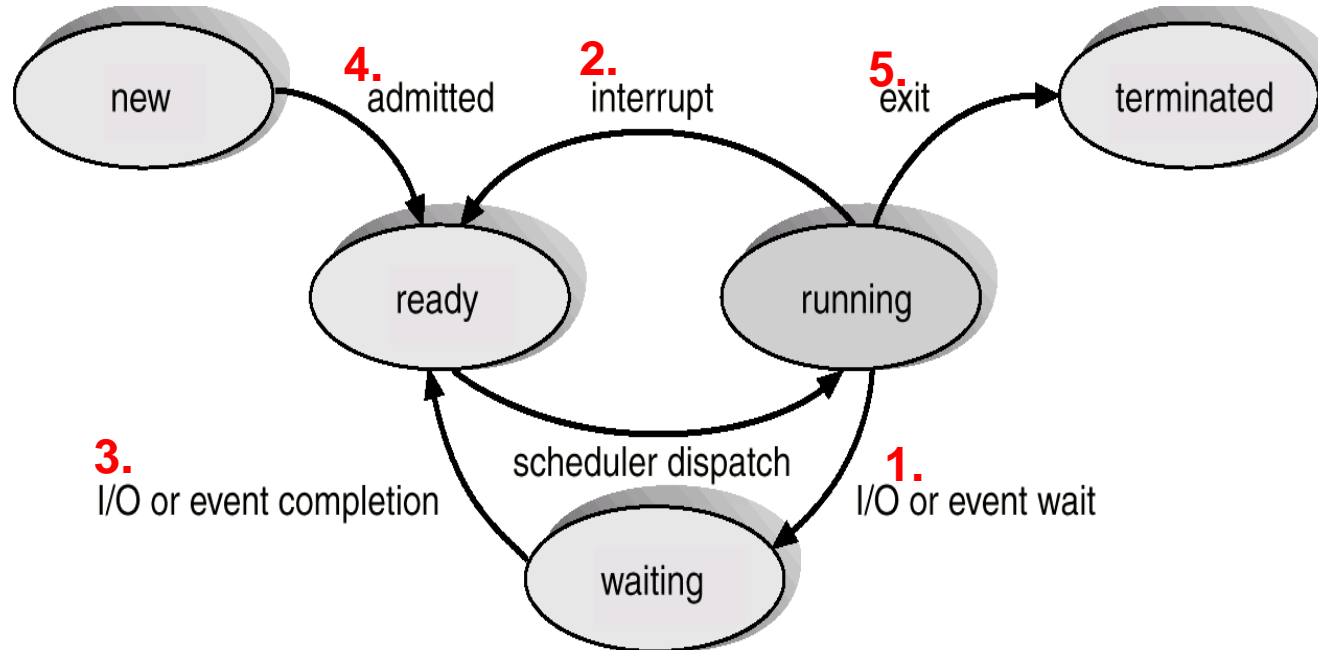
- CPU–I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait.
 - CPU burst: the period for CPU executions.
 - I/O burst: the period for I/O wait.
- The goal of scheduling is to keep the CPU busy (i.e., improve CPU utilization with multiprogramming).
- For simplicity, we focus on the concepts of **short-term scheduler** on a single CPU.

Alternating Sequence of CPU And I/O Bursts



CPU Scheduler

- Selects one of the processes in the ready queue, and allocates the CPU to it.
- CPU scheduling decisions may take place when a process changes state.
 - Which case(s) does scheduling **must** occur in order to keep the CPU busy?



CPU Scheduler (cont.)

- For situations 1 and 5, there is no choice.
However, there is a choice for situations 2, 3 & 4.
 - **Nonpreemptive** means that once the CPU has been allocated to a process, the process keeps the CPU until it releases the CPU either by terminating or requesting I/O/event wait.
 - * *Scheduling is nonpreemptive if it happens **only** under situations 1 and 5*
 - **Preemptive** means that the CPU can be taken away from running process at any time by OS.
 - * *Scheduling is preemptive if it **also** happens under situations 2, 3 and 4*

Scheduling Criteria

- CPU utilization – keep the CPU as busy as possible
- Throughput – # of processes that complete their execution per time unit
- Turnaround time – amount of time to execute a particular process, i.e. from the time of submission to the time of completion.

Scheduling Criteria

- Waiting time – amount of time a process has been waiting in the ready queue.
 - Turn-around time consists of three components: CPU burst, I/O burst, and waiting time.
 - If all processes have a single CPU burst (no I/O),
waiting time = turn-around time – CPU burst time.
- Response time – amount of time it takes from when a request was submitted until the **first** response is produced

Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- **Min waiting time**
- Min response time

Overview of Scheduling Algorithms

- First-Come, First-Served (FCFS)
- Shortest Job First (SJF)
- Priority Scheduling
- Round Robin (RR)
- Multilevel Queue Scheduling

First-Come, First-Served (FCFS) Scheduling

- Example:

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that processes arrive in the order: P_1 , P_2 , P_3 all at time 0.
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1.$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case.

FCFS Scheduling (Cont.)

- Problem of FCFS:
 - will cause *convoy effect*, that means short processes suffer because they have to wait for long process to finish.


Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time.
- Two schemes:
 - nonpreemptive – once CPU given to the process it cannot be preempted until completes its CPU burst.
 - Preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (**SRTF**).
- SJF is optimal – gives minimum average waiting time for a given set of processes.

Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7 ✓
P_2	2.0	4 ✓
P_3	4.0	1 ✓
P_4	5.0	4 ✓

- SJF (non-preemptive)




- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Example of Preemptive SJF (i.e. SRTF)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7 \rightarrow 5 \rightarrow 0
P_2	2.0	4 \rightarrow 2 \rightarrow 0
P_3	4.0	1 \rightarrow 0
P_4	5.0	4 \rightarrow 0

- SJF (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority).
 - Preemptive
 - nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time.
- Problem \equiv Starvation – low priority processes may never execute.
- Solution \equiv **Aging** – as time progresses increase the priority of the process.


Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then no process waits more than $(n-1)q$ time units.
- Performance
 - q large \Rightarrow FCFS
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high.

Example: RR with Time Quantum = 20


<u>Process</u>	<u>Burst Time</u>
P_1	53 \rightarrow 33
P_2	17 \rightarrow 0
P_3	68 \rightarrow 48
P_4	24 \rightarrow 4

- The Gantt chart is (all processes come at time zero, in the order of P_1 , P_2 , P_3 , P_4):




- Typically, higher average turnaround than SJF, but better *response time*.

How a Smaller Time Quantum Increases Context Switches



Turnaround Time Varies With The Time Quantum



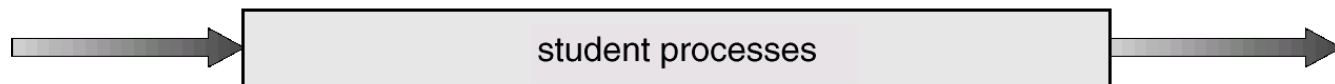



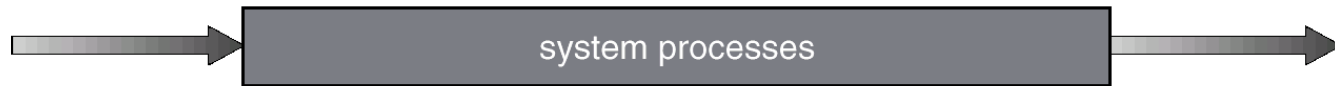
process	time
P_1	6
P_2	3
P_3	1
P_4	7

Multilevel Queue

- Ready queue is partitioned into separate queues:
foreground (e.g. interactive)
background (e.g. batch)
- Each queue has its own scheduling algorithm,
foreground – e.g. RR
background – e.g. FCFS

Multilevel Queue Scheduling

highest priority



lowest priority

Multilevel Queue (Cont.)

- Scheduling must be done between the queues.
 - Fixed priority scheduling; i.e., serve all from foreground then from background. Possibility of **starvation** for the background processes.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; examples are:
 - * 80% to foreground queue
 - * 20% to background queue

Advanced Readings

- The following animations are very good tools for learning CPU scheduling.
 - <http://www.utdallas.edu/~ilyen/animation/cpu/program/prog.html>
 - <http://cs.uttyler.edu/Faculty/Rainwater/COSC3355/Animations/>