

Part 5: Deadlocks and Starvation


- Deadlock Problem
- System Model
- Deadlock Conditions
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection

The Deadlock Problem

- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
- Example
 - System has 2 different type of tape drives; **both P1 and P2 need the two tapes to finish execution.**
 - P_1 and P_2 each hold one tape drive and each needs another one.
- Example
 - semaphores A and B , initialized to 1

| | |
|----------------------|----------------------|
| P_0 | P_1 |
| <i>wait</i> (A); | <i>wait</i> (B); |
| <i>wait</i> (B); | <i>wait</i> (A); |

Bridge Crossing Example



- Allow only one car on the bridge.
- The bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (release resource, i.e. the bridge, and rollback).
- Several cars may have to be backed up if a deadlock occurs.

System Model

- Resource types R_1, R_2, \dots, R_m
e.g. memory space, I/O devices
- Each resource type R_i has W_i identical instances.
- Each process utilizes a resource as follows:
 - request
 - use
 - release


Resource-Allocation Graph

A set of vertices V and a set of edges E .

- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.
- request edge: directed edge $P_i \rightarrow R_j$
 - *When the request is granted, the request edge is removed.*
- assignment edge: directed edge $R_j \rightarrow P_i$
 - *When the resource is released, the assignment edge is removed.*

Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances




- P_i requests instance of R_j



- P_i is holding an instance of R_j




Example of a Resource Allocation Graph




No deadlock.

$P_3 \rightarrow P_2 \rightarrow P_1$

Resource Allocation Graph With A Deadlock



Resource Allocation Graph With A Cycle But No Deadlock



Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, **possibility of** deadlock.

Deadlock Conditions

Deadlock **can** arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource.
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.

Deadlock Conditions (Cont.)

- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for P_2 , ..., P_{n-1} is waiting for P_n , and P_n is waiting for P_0 .

Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state.
- Allow the system to enter a deadlock state and then recover.
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX.



Deadlock Prevention

Prevent at least one of the deadlock conditions to happen.

- Mutual Exclusion – not required for sharable resources; must hold for nonsharable resources.
 - Sharable resources: code section, read-only data areas.
 - Nonshareable resources: printers, data areas to be written, chopstick in Dining-Philosophers.
- Hold and Wait – must guarantee that whenever a process requests a resource, it does not hold any other resources.
 - In Dining-Philosophers, allow a philosopher to pick up his chopsticks only if both chopsticks are available.

Deadlock Prevention

Prevent at least one of the deadlock conditions to happen.

- No Preemption – If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released. The process restarts later.
 - In Dining-Philosophers, if a philosopher cannot get another chopstick for a long time, preempt.
- Circular Wait –
 - In Dining-Philosophers, allow at most **four** philosophers to be hungry simultaneously.
 - Impose a total ordering of all resource types (see backup slides).

Deadlock Avoidance

- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that the system never goes into **unsafe** state.
 - When a process requests an available resource, system must decide if immediate allocation leaves the system in a **safe** state.
 - If safe, the request is granted. Otherwise, the process must wait.
- System is in safe state if there exists a *safe sequence* of all processes.

Safe State (Cont.)

- Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is **safe** if for each P_i , the resources that P_i requests can be satisfied by currently available resources + resources held by all the P_j , with $j < i$.
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished.
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate.
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on. → All processes in the sequence can finish.

Example of Safe Sequence

Available: 1

| Processes | Hold | Request |
|-----------|------|---------|
| P1 | 1 | 1 |
| P2 | 1 | 2 |
| P3 | 1 | 3 |

Q1: $\langle P1, P2, P3 \rangle$ safe?

Yes.

$P1.request \leq Available$

$P2.request \leq P1.Hold + Available$

$P3.request \leq P1.Hold + P2.Hold + Available$

Q2: $\langle P3, P2, P1 \rangle$ safe?

No.

$P3.request > Available.$

Q3: Is the system safe?


Yes.

$\langle P1, P2, P3 \rangle$ is the safe seq.

Basic Facts

- If a system is in safe state \Rightarrow no deadlocks.
- If a system is in unsafe state \Rightarrow possibility of deadlock.
 - Consider the process can release resources before termination.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.

Safe, unsafe, deadlock state spaces



Banker's Algorithm

- An algorithm decides whether to allocate the resource and to avoid the unsafe state.
- For multiple instances of each resource type.
- Each process must declare max. # of instances of each resource type that it needs.
- When a process requests a resource it may have to wait (if the allocation leads to unsafe state).
- When a process gets all its resources it must return them in a finite amount of time.

Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- *Available*: Vector of length m . If **Available** $[j] == k$, there are k instances of resource type R_j available.
- *Max*: $n \times m$ matrix. If $Max [i,j] == k$, then process P_i may request at most k instances of resource type R_j .
- *Allocation*: $n \times m$ matrix. If $Allocation[i,j] == k$ then P_i is currently allocated k instances of R_j .
- *Need*: $n \times m$ matrix. If $Need[i,j] == k$, then P_i may need k more instances of R_j to complete its task.

$$Need [i,j] = Max[i,j] - Allocation [i,j].$$

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

Work = *Available*

Finish [*i*] = *false* for *i* = 1, 2, ..., *n*.

2. Find an *i* such that both:

(a) *Finish* [*i*] == *false*

(b) $Need_i \leq Work$

What if multiple answers?

If no such *i* exists, go to step 4.

3. *Work* = *Work* + *Allocation*_{*i*}

Finish[*i*] = *true*

go to step 2.

Work = *Available* + resources held by all the *P_j* (*Finish*[*j*] = *true*).

4. If *Finish* [*i*] == *true* for all *i*, then the system is in a safe state. Otherwise, the system is unsafe.

Example of Banker's Algorithm

- 5 processes P_0 through P_4 ; 3 resource types: A (10 instances), B (5 instances) and C (7 instances).
- Snapshot at time T_0 :

| | <u>Allocation</u> | <u>Max</u> | <u>Available</u> |
|-------|-------------------|------------|------------------|
| | A B C | A B C | A B C |
| P_0 | 0 1 0 | 7 5 3 | 3 3 2 |
| P_1 | 2 0 0 | 3 2 2 | |
| P_2 | 3 0 2 | 9 0 2 | |
| P_3 | 2 1 1 | 2 2 2 | |
| P_4 | 0 0 2 | 4 3 3 | |

Example of Banker's Algorithm (Cont.)

- Then the matrix Need is defined to be Max – Allocation.

| <u>Allocation</u> | <u>Max</u> | <u>Need</u> | <u>Available</u> |
|-------------------|------------|-------------|------------------|
| A B C | A B C | A B C | A B C |
| P_0 0 1 0 | 7 5 3 | 7 4 3 | 3 3 2 |
| P_1 2 0 0 | 3 2 2 | 1 2 2 | |
| P_2 3 0 2 | 9 0 2 | 6 0 0 | |
| P_3 2 1 1 | 2 2 2 | 0 1 1 | |
| P_4 0 0 2 | 4 3 3 | 4 3 1 | |

Example of Banker's Algorithm (Cont.)

- The working is as follows:

| <u>Need</u> | | <u>Allocation</u> | | | Work | | | Finished Process | |
|-------------|---|-------------------|---|---|-------------|---|---|------------------|-------|
| | A | B | C | | A | B | C | | |
| P_0 | 7 | 4 | 3 | ✓ | P_0 | 0 | 1 | 0 | |
| + | | | | | 3 | 3 | 2 | | |
| P_1 | 1 | 2 | 2 | ✓ | P_1 | 2 | 0 | 0 | P_1 |
| + | | | | | 5 | 3 | 2 | | |
| P_2 | 6 | 0 | 0 | ✓ | P_2 | 2 | 1 | 1 | P_3 |
| + | | | | | 7 | 4 | 3 | | |
| P_3 | 0 | 1 | 1 | ✓ | P_3 | 0 | 0 | 2 | P_4 |
| + | | | | | 7 | 4 | 5 | | |
| P_4 | 4 | 3 | 1 | ✓ | P_4 | 3 | 0 | 2 | P_2 |
| + | | | | | 10 | 4 | 7 | | |
| + | | | | | 0 | 1 | 0 | P_0 | |
| | | | | | 10 | 5 | 7 | | |

Example of Banker's Algorithm (Cont.)

- The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria.
- There are multiple safe sequences for that system state.

Resource-Request Algorithm for Process P_i

$Request_i$ = request vector for process P_i . If $Request_i[j] == k$ then process P_i wants k instances of resource type R_j .

$V1 \leq V2 \iff V1[j] \leq V2[j], \text{ for all } j$

The i^{th} row of the matrix.

1. If $Request_i \leq Need_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available.

Resource-Request Algorithm for Process P_i (Cont.)

3. **Pretend** to allocate requested resources to P_i by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i.$$

4. Then run the safety algorithm, if the result is:

- safe \Rightarrow the resources are allocated to P_i .
- unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Example of Banker's Algorithm (Cont.)

- Suppose P_1 requests (1, 0, 2) resources.
- Step1: Check $\text{Request}_1 \leq \text{Need}_1$, $(1,0,2) \leq (1,2,2) \Rightarrow \text{true}$
- Step2: Check that $\text{Request}_1 \leq \text{Available}$, i.e. $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$.
- Steps 3 & 4:

Allocation Max Need Available

| | A | B | C | A | B | C | A | B | C | A | B | C |
|-------|--------------|--------------|--------------|---|---|---|--------------|--------------|--------------|--------------|--------------|--------------|
| P_0 | 0 | 1 | 0 | 7 | 5 | 3 | 7 | 4 | 3 | 3 | 3 | 2 |
| P_1 | 2 | 0 | 0 | 3 | 2 | 2 | 1 | 2 | 2 | 2 | 3 | 0 |
| | 3 | 0 | 2 | | | | 0 | 2 | 0 | | | |
| P_2 | 3 | 0 | 2 | 9 | 0 | 2 | 6 | 0 | 0 | | | |
| P_3 | 2 | 1 | 1 | 2 | 2 | 2 | 0 | 1 | 1 | | | |
| P_4 | 0 | 0 | 2 | 4 | 3 | 3 | 4 | 3 | 1 | | | |

**Pretend to
allocate
requested
resources**

Example of Banker's Algorithm (Cont.)

Executing Safety Algorithm

| | <u>Allocation</u> | <u>Need</u> | <u>Available</u> |
|-------|-------------------|------------------|------------------|
| | A B C | A B C | A B C |
| P_0 | 0 1 0 | 7 4 3 | 3 3 2 |
| P_1 | 2 0 0 | 1 2 2 | 2 3 0 |
| P_2 | 3 0 2 | 6 0 0 | |
| P_3 | 2 1 1 | 0 1 1 | |
| P_4 | 0 0 2 | 4 3 1 | |

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement.

Example of Banker's Algorithm (Cont.)

- Can further request for (1,0,2) by P_1 be granted?
 - Check if $\text{Request}_1 \leq \text{Need}_1$, $(1,0,2) \not\leq (0,2,0) \Rightarrow \text{false} \Rightarrow \text{error}$
- Can further request for (0,2,0) by P_0 be granted?
 - Check if $\text{Request}_0 \leq \text{Need}_0$, $(0,2,0) \leq (7,4,3) \Rightarrow \text{true}$
 - Check that $\text{Request}_0 \leq \text{Available}$, i.e. $(0,2,0) \leq (2,3,0) \Rightarrow \text{true}$.

Allocation Need Available

A B C A B C A B C

P_0 ~~0 1 0~~ ~~7 4 3~~ ~~2 3 0~~
 0 3 0 7 2 3 2 1 0

P_1 3 0 2 0 2 0

P_2 3 0 2 6 0 0

P_3 2 1 1 0 1 1

P_4 0 0 2 4 3 1

☹ We cannot find $\text{Need}_i < \text{Available}$
 → Restore the state

- How about a request for (2,3,0) by P_4 ?

Comments are Welcome

- Online feedback systems welcome your feedbacks (Due)
- Your comments are
 - To have a self-a
 - To improve the future teaching
- Speak out:
 - If you love my lecture, please treasure the good
 - If you think of an idea for my next lecture...



THANK
YOU!



GOOD
LUCK!

Deadlock Detection

- Allow system to enter deadlock state
- Then invoke detection algorithms. There are two:
 - For single instance of each resource type (refer to textbook)
 - For multiple instances of each resource type
- Then invoke recovery algorithm (refer to textbook)

This slide and the later slides in this chapter are not examinable.

Multiple Instances of Each Resource Type

- **Available:** A vector of length m indicates the number of available resource types.
- **Allocation:** An $n \times m$ matrix defines the number of resources of each type currently allocated to each process.
- **Request:** An $n \times m$ matrix indicates the current request of each process. If $Request[i][j] = k$, then process P_i is requesting k more instances of resource type R_j .

Detection Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:
 - (a) *Work* = *Available*
 - (b) For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then $Finish[i] = \text{false}$; otherwise, $Finish[i] = \text{true}$.
2. Find an index *i* such that both:
 - (a) $Finish[i] == \text{false}$
 - (b) $Request_i \leq Work$If no such *i* exists, go to step 4.

Detection Algorithm (Cont.)

3. $Work = Work + Allocation_i$

$Finish[i] = true$

go to step 2.

4. If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if $Finish[i] == false$, then P_i is deadlocked.

Algorithm requires an order of $m \times n^2$ operations to detect whether the system is in deadlocked state.

Example of Detection Algorithm

- Five processes P_0 through P_4 ; three resource types A (7 instances), B (2 instances), and C (6 instances).
- Snapshot at time T_0 :

| | <u>Allocation</u> | | | <u>Request</u> | | | <u>Available</u> | | |
|-------|-------------------|---|---|----------------|---|---|------------------|---|---|
| | A | B | C | A | B | C | A | B | C |
| P_0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| P_1 | 2 | 0 | 0 | 2 | 0 | 2 | | | |
| P_2 | 3 | 0 | 3 | 0 | 0 | 0 | | | |
| P_3 | 2 | 1 | 1 | 1 | 0 | 0 | | | |
| P_4 | 0 | 0 | 2 | 0 | 0 | 2 | | | |

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $Finish[i] = \text{true}$ for all i .

Example of Detection Algorithm(Cont.)

- Suppose P_2 requests an additional instance of type C.

| <u>Allocation</u> | <u>Request</u> | <u>Available</u> |
|-------------------|----------------|------------------|
| A B C | A B C | A B C |
| P_0 0 1 0 | 0 0 0 ✓ | 0 0 0 |
| P_1 2 0 0 | 2 0 2 ☹️ | |
| P_2 3 0 3 | 0 0 1 ☹️ | |
| P_3 2 1 1 | 1 0 0 ☹️ | |
| P_4 0 0 2 | 0 0 2 ☹️ | |

| Work | Finished |
|---------|----------|
| A B C | P_0 |
| 0 0 0 | |
| + 0 1 0 | |
| 0 1 0 | |


Example of Detection Algorithm(Cont.)

- State of system?
 - Can reclaim resources held by process P_0 , but insufficient resources to fulfill other processes requests.
 - Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4 .


Advanced Readings

- “The Deadlock Problem: An Overview”, By Sreekaanth S. Isloor, T. Anthony Marsland. (pdf in EdveNTUre)
- Other readings
 - Java Concurrency,
<http://docs.oracle.com/javase/tutorial/essential/concurrency/index.html>
 - Deadlock,
<http://en.wikipedia.org/wiki/Deadlock>

An Example of Preventing Circular Wait



An Example of Preventing Circular Wait (Cont')



Consider the worst case
if A gets only chopstick 0.