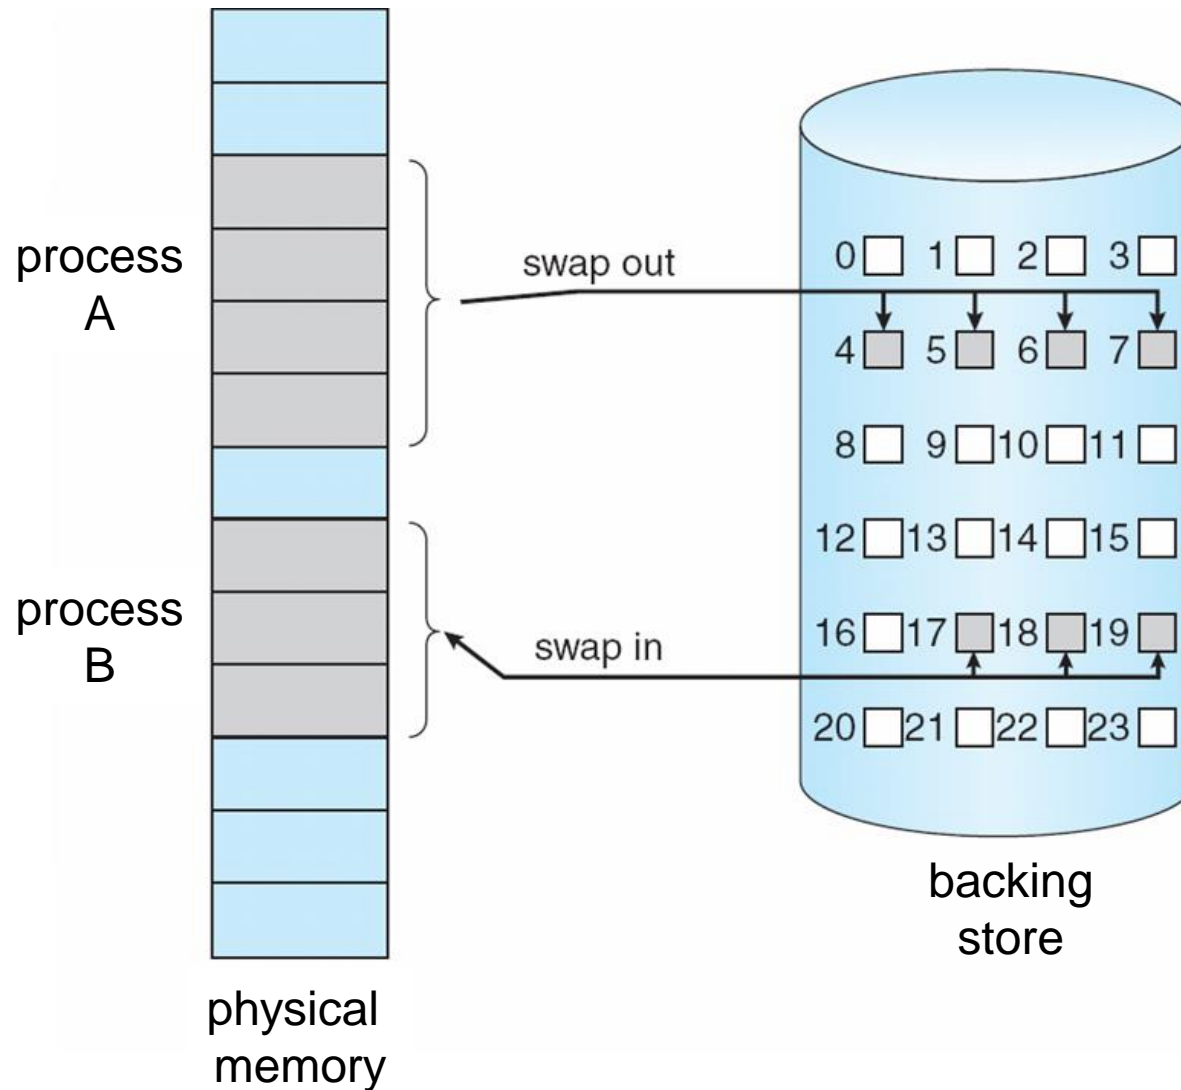# Part 7: Virtual Memory

- Introduction
  - Swapping, Virtual Memory Support

- Demand Paging
  - Page Fault, Performance of Demand Paging, Page Replacement

- Page-Replacement Algorithms
  - FIFO, Optimal, LRU, LRU Implementation, and LRU Approximation

- Allocation of Frames
  - Fixed vs. Variable Allocation, Local vs. Global Replacement

- Thrashing
  - Cause of Thrashing, Working-Set Model, Page-fault Frequency Scheme

Introduction - Swapping

- A process can be *swapped* temporarily out of memory to a *backing store*, and then brought back into memory for continued execution

- Backing store – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images

- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped

- Swapping can be used to
  - Make space available for processes that require more memory
  - Increase the *degree of multiprogramming (i.e., medium term scheduler)*
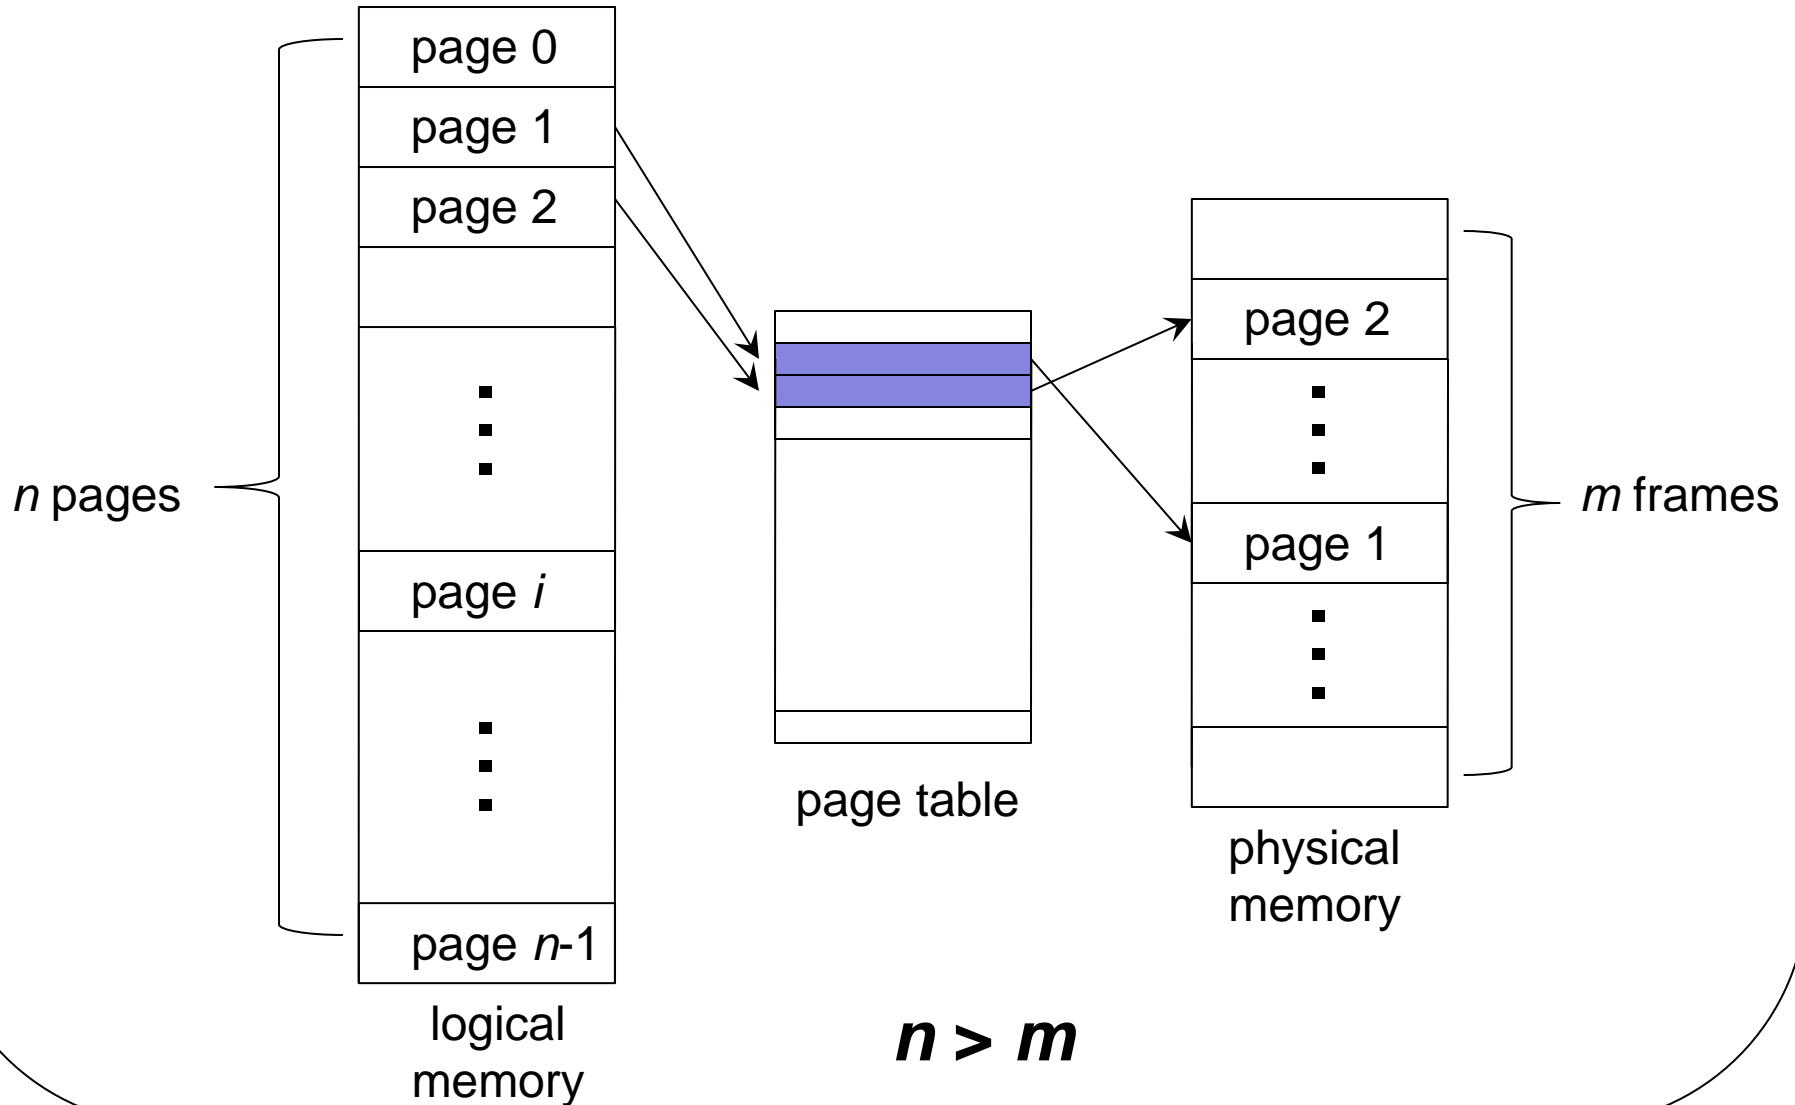
Introduction – Swapping (cont.)

process A

process B

physical memory

swap out

swap in

backing store

0 □ 1 □ 2 □ 3 □
4 ■ 5 ■ 6 ■ 7 ■
8 □ 9 □ 10 □ 11 □
12 □ 13 □ 14 □ 15 □
16 □ 17 ■ 18 ■ 19 ■
20 □ 21 □ 22 □ 23 □

Introduction – Virtual Memory Support

- So far we assumed that the whole program is loaded into the memory, at least at some point of time.
  Restriction: Overall program size, including all library routines that may be linked at runtime, must be restricted to the size of main memory

- To remove this restriction, the logical memory size requirement of the user program must be de-coupled from the size of the physical memory space

- This is what virtual memory aims to achieve

  – The size of virtual memory is limited by the address scheme of a computer, not by the actual size of the physical memory.

7.4

Introduction – Virtual Memory Support (Cont.)

page 0
page 1
page 2

page *i*

page *n*-1

n pages

logical memory

page table

page 2

page 1

physical memory

m frames

n > m

Introduction – Virtual Memory Support (cont.)

- Virtual Memory Support: separation of user logical memory from the physical memory

- Implications of Virtual Memory Support:

  - Logical address space can be much larger than physical address space

  - Only part of the program needs to be in memory for execution $\Rightarrow$ greater degree of multiprogramming $\Rightarrow$ increased CPU utilization

  - Need to allow pages to be paged in and out

- Virtual memory is implemented (mostly) using *demand paging* (demand segmentation is also used)

Demand Paging

- Memory requirement of a process is divided into pages

- Page is needed $\Rightarrow$ page reference is generated $\Rightarrow$ logical address <page no., offset>

- With each page table entry, a *valid–invalid bit* is associated

    - **v** $\Rightarrow$ in-memory, page table entry contains location of the page in memory (i.e., frame number)

    - **i** $\Rightarrow$ not-in-memory, page table entry contains the address of the page on disk

    - Initially valid–invalid bit is set to **i** on all entries

    - During address translation, if valid–invalid bit in page table entry is **i** $\Rightarrow$ *page fault* (to bring the page into memory)

Demand Paging (Cont.)

generate a reference to page 6

PC

Page Fault

load G

page is on backing store ③

operating system

reference

trap ②

①

restart instruction ⑥

page table · i

k

free frame

reset page table ⑤

bring in missing page ④

(5.1) Modify the valid-invalid bit into valid
(5.2) Set the page table entry to be k

physical memory
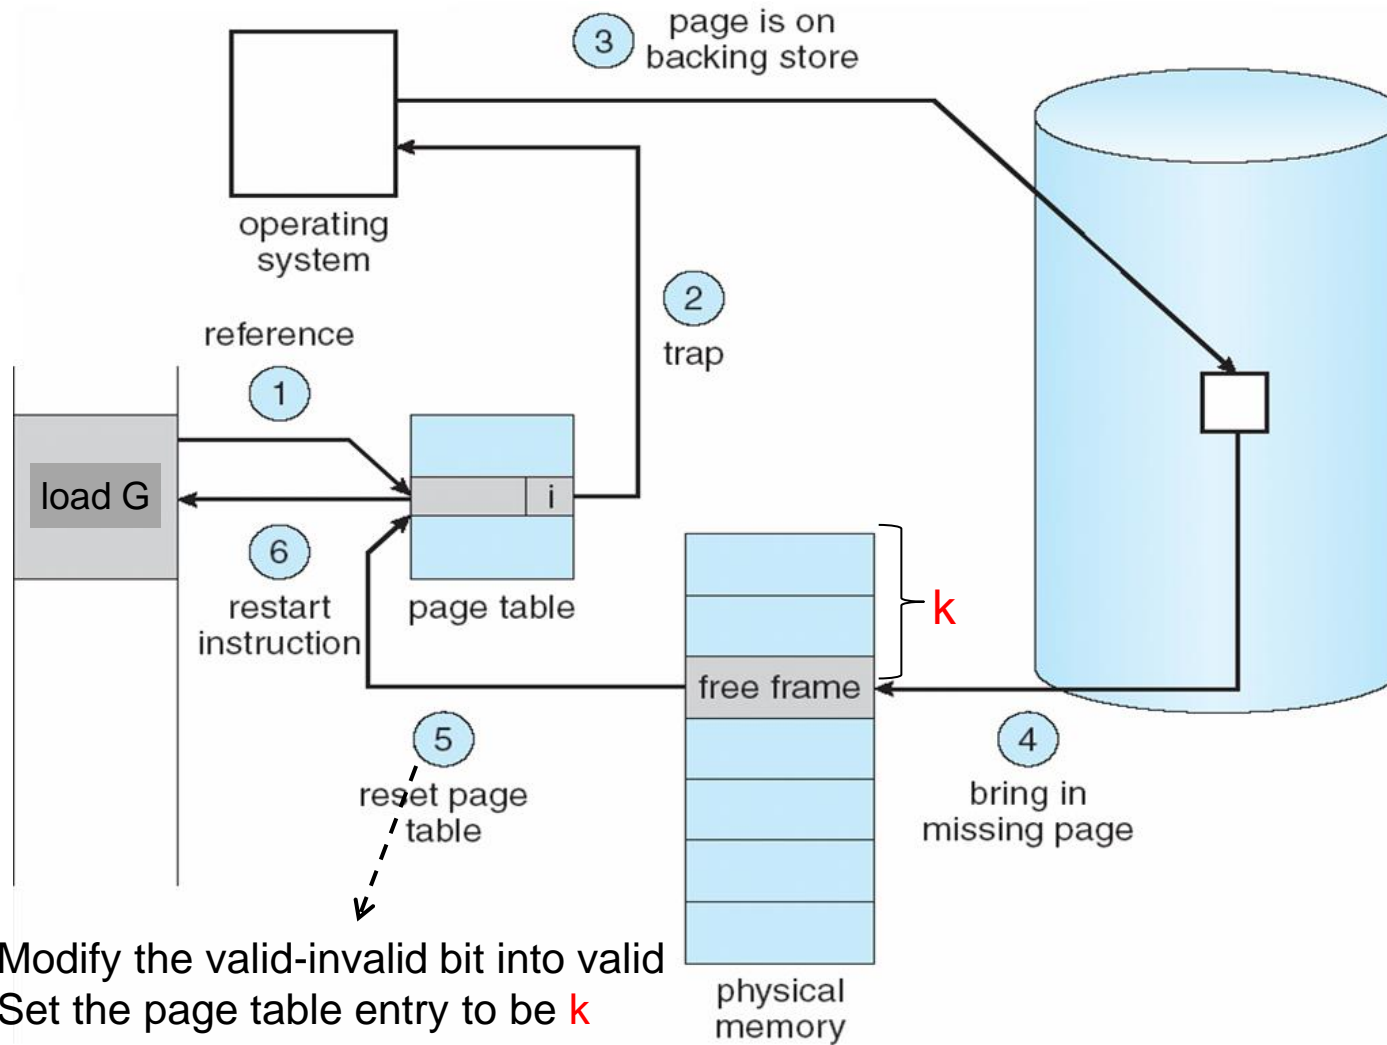
Page Fault (Cont.)

1. If there is a reference to a page, page table is checked to determine if the page is in the memory.

2. If the page is not in the memory, trap to OS to bring in the page. The state of interrupted process is saved.

3. OS uses the page table entry to locate the page in backing store.

4. OS finds a free frame (i.e., using free frame list & page replacement algorithm), and schedules a disk operation to read the desired page into the newly allocated frame.

5. When the disk read is complete, page table entry is modified.

6. The state of interrupted process is resumed, and the instruction is now re-started.

Performance of Demand Paging

- Let p = probability of Page Fault, $0 \leq p \leq 1.).0$
  - if $p = 0$ no page faults
  - if $p = 1$, every reference causes a fault

- Effective Access Time (EAT)

$$EAT = (1 - p) \text{ x memory access time}$$

$$+ p \text{ x page fault time}$$

- In practice, memory access time << page fault time.

Performance of Demand Paging (Cont.)

Three major components of page fault time:

- Servicing the page fault interrupt

- Read in a new page (longest component)

- Restart the process

Page Replacement

- When page fault occurs, what happens if there is no free frame?
  Solution: Find some page in memory, not really in use, and page it out for only dirty page (i.e., replace page)

- If no free frame is available, then two page transfers are necessary, one out (for dirty page only) and one in. This increases further the page fault time!

- Use a *page replacement algorithm* to locate the page to be replaced. Desirably, the algorithm should result in minimum number of page faults

Page Replacement (Cont.)

generate a reference to page 3

not in the memory

memory full

logical memory for user 1

page table for user 1

logical memory for user 2

page table for user 2

physical memory

valid–invalid bit

frame

Page Replacement (Cont.)

frame   valid–invalid bit

0	i
f	v

page table

② change to invalid

④ reset page table for new page

f | victim

physical memory

① page out

③ page in

1. Find a *victim* page using place replacement algorithm and write the victim page to the backing store (*page-out for dirty page*)

2. Change the page table entry of the victim page to invalid

   Why?

3. Bring the desired page into the newly freed frame (*page-in*)

4. Update the page table entry for the new page

Page Replacement (Cont.)

- This overhead can be minimized by not writing pages that have not been modified since they were brought into memory. Such pages are identified by using a *modify* (or *dirty*) bit

- When replacing pages, **only** *dirty pages* are actually written to backing store

- Page replacement completes separation between logical memory and physical memory - large virtual memory can be provided on a smaller physical memory

Page-Replacement Algorithms

- Objective is to have an algorithm with lowest page-fault rate

- Evaluate algorithm by running it on a particular string of memory references (reference string), counting the number of page faults

- In all our examples, the ref. string used is

$$1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5$$

Page-Replacement Algorithms (Cont.)

- Obviously, as the # of available frames increases, the # of page faults will decrease.

- However, there are exceptions!

7.18

Outline

- First-In-First-Out (FIFO) Algorithm

- Optimal Algorithm

- Least Recently Used (LRU) Algorithm

- Second-Chance Algorithm

- Clock Algorithm

First-In-First-Out (FIFO) Algorithm

- The easiest and simplest algorithm.

  Replace the page that has been *loaded* in the memory for the longest period of time

- Example using 3 free frames:

		1	2	3	4	1	2	5	1	2	3	4	5
F1		1	1	1	4	4	4	5	5	5	5	5	5
F2			2	2	2	1	1	1	1	1	3	3	3
F3				3	3	3	2	2	2	2	2	4	4
		P	P	P	P	P	P	P			P	P	

9 page faults

FIFO Algorithm (Cont.)

– Example using 4 free frames:

		1	2	3	4	1	2	5	1	2	3	4	5
F1		1	1	1	1	1	1	5	5	5	5	4	4
F2			2	2	2	2	2	2	1	1	1	1	5
F3				3	3	3	3	3	3	2	2	2	2
F4					4	4	4	4	4	4	3	3	3
		P	P	P	P			P	P	P	P	P	P

10 page faults !

- Belady's Anomaly

– more frames $\not\Rightarrow$ fewer page faults

Optimal Algorithm

- Replace the page that will not be used for the longest period of time.

		1	2	3	4	1	2	5	1	2	3	4	5
F1		1	1	1	1	1	1	1	1	1	1	4	4
F2			2	2	2	2	2	2	2	2	2	2	2
F3				3	3	3	3	3	3	3	3	3	3
F4					4	4	4	5	5	5	5	5	5
		P	P	P	P			P				P	

6 page faults

- *Problem: How will the OS know the future references*?

Optimal Algorithm (Cont.)

- Theoretically best but practically infeasible

- Used as a benchmark to measure how well other algorithms perform

- Does not suffer from Belady's anomaly:

  Add a frame $\Rightarrow$ more look-ahead $\Rightarrow$

  Inclusion property : Pages loaded in n frames is always a subset of pages in $n+1$ frames

Least Recently Used (LRU) Algorithm

- Replace the page that has not been used for the longest period of time

		1	2	3	4	1	2	5	1	2	3	4	5
F1		1	1	1	1	1	1	1	1	1	1	1	5
F2			2	2	2	2	2	2	2	2	2	2	2
F3				3	3	3	3	5	5	5	5	4	4
F4					4	4	4	4	4	4	3	3	3
		P	P	P	P			P			P	P	P

8 page faults

- Similar to Optimal Algorithm, but looking backward in time (and not susceptible to Belady's anomaly)

Implementation of LRU Algorithm

- Counter Implementation

  – Each page table entry has a time-of-use field

  – The CPU increments a logical clock for every memory reference

  – Whenever a page is referenced, the value of the clock is copied to the time-of-use field of the page

  – When a page fault occurs, the page with the smallest time-of-use value is kicked out

7.25

Implementation of LRU Algorithm (Cont.)

- Stack implementation

  - Keep a stack of page numbers in a doubly linked list

  - When a page is referenced:

    ➢ move it to the top

    ➢ the bottom of the stack is the LRU page

  - List update is expensive, but no search is needed for page replacement

Implementation of LRU Algorithm (Cont.)

- Stack Implementation Example

1	2	3	4	1	2	5	1	2	3	4	5
1	2	3	4	1	2	5	1	2	3	4	5
	1	2	3	4	1	2	5	1	2	3	4
		1	2	3	4	1	2	5	1	2	3
			1	2	3	4	4	4	5	1	2
						3			4	5	1
P	P	P	P			P			P	P	P

8 Page Faults

Stack of Page Numbers

LRU Approximation Algorithms

- Hardware support necessary for exact algorithms are expensive and generally not available

- However, many systems support a hardware settable *reference bit*

  - With each page associate a bit R, initially = 0
  - When page is referenced, hardware sets its R to 1
  - Replace page with $R = 0$ (if one exists). Note that the exact order of use cannot be determined by R
    - ❄ There may be many pages with R = 1

- Approximation Algorithm:

  - Second-Chance (Clock) algorithm

Second-Chance Algorithm

- A variation of FIFO algorithm

- Choose oldest page as candidate to be replaced

- If the reference bit R of the candidate is set (R=1), the page is given a second chance:
  - the R bit is cleared and the page is treated as if it has just been read in

- Otherwise, the page is kicked out

7.29

Clock Algorithm

- Implementation of the Second-Chance algorithm

- The clock hand points to the oldest page

- When a page fault occurs, the page being pointed by the hand is inspected

- If R bit is set, it is cleared and the hand is advanced to the next page

- Otherwise, the page is evicted and new page is inserted into its place. The hand is advanced one position

- In the worst case, when all the bits are set, the hand cycles through the whole queue, giving each page a second chance. In this case it degenerates to FIFO replacement

Clock Algorithm (Cont.)

newer

7.31

Example: Clock Algorithm

1	2	3	4	1	2	5	1	2	3	4	5
1,0	2,0	3,0	4,0	4,0	4,0	4,0	4,0	4,0	3,0	3,0	5,0
	1,0	2,0	3,0	3,0	3,0	5,0	5,0	5,0	5,0	4,0	4,0
		1,0	2,0	2,0	2,1	2,0	2,0	2,1	2,1	2,0	2,0
			1,0	1,1	1,1	1,0	1,1	1,1	1,1	1,0	1,0
						3			4	5	3
P	P	P	P			P			P	P	P

8 Page Faults

Circular List of Page Numbers

The second digit of each entry is the *R* bit. The shaded entry is where the clock hand points to. It moves upwards. When reaching the top, it will go back to the bottom

7.32

Virtual Memory

Allocation of Frames

- The OS must decide how many pages to bring into main memory

  - the smaller the amount of memory allocated to each process, the more processes can reside in memory

  - small number of pages loaded increases page faults

  - beyond a certain size, further allocations of pages will not effect the page fault rate

Fixed vs. Variable Allocation

- *Fixed Allocation*
  - gives a process a fixed number of frames in main memory within which to execute
    - ❆ e.g., equal allocation or proportional allocation
  - when a page fault occurs, one of the pages of that process must be replaced

- *Variable Allocation*
  - allows the number of page frames allocated to a process to be varied over the lifetime of the process

7.34

Global vs. Local Allocation

- *Global replacement* – process selects a replacement frame from the set of all frames; one process can take a frame from another process

    $\Rightarrow$ Performance depends on external circumstances (other processes)

- *Local replacement* – each process selects only from its own set of allocated frame

    $\Rightarrow$ may hinder other processes by not making available its less used pages/frames

7.35

Global vs. Local Allocation (Cont.)

Time Last Accessed

Process A generates a reference to page A6

	Time Last Accessed
A0	10
A1	7
A2	5
A3	4
A4	6
A5	3
B0	9
B1	4
B2	6
B3	2
B4	5
B5	6
B6	12
C1	3
C2	5
C3	6

original configuration

A0
A1
A2
A3
A4
(A6)
B0
B1
B2
B3
B4
B5
B6
C1
C2
C3

local replacement

A0
A1
A2
A3
A4
A5
B0
B1
B2
(A6)
B4
B5
B6
C1
C2
C3

global replacement

7.36

Allocation of Frames – Summary

	Local Replacement	Global Replacement
Fixed Allocation	•Number of frames allocated to a process is fixed. •Page to be replaced is chosen from among the frames allocated to that process.	•Not possible.
Variable Allocation	•The number of frames allocated to a process may be changed from time to time to maintain the working set of the process. •Page to be replaced is chosen from among the frames allocated to that process.	•Page to be replaced is chosen from all available frames in main memory; this causes the size of the resident set of processes to vary.

Thrashing

- If a process does not have "enough" pages, the page-fault rate is very high. This leads to:
    - Low CPU utilization
    - OS thinks that it needs to increase the degree of multiprogramming
    - Another process is added to the system
    - CPU utilization drops further

7.38

Thrashing (Cont.)

- *Thrashing* $\equiv$ a process is busy bringing pages in and out (no work is being done)

7.39

Strategies to combat Thrashing

- Two approaches are used:
    - Working-Set Model
    - Page-Fault Frequency

Working-Set Model

- *Locality of Reference*: Processes tend to refer to pages in a localised manner

    – *Temporal Locality*: locations referred to just before are likely to be referred to again (i.e., location references are clustered over time)

    – *Spatial Locality*: code & data are usually clustered physically

- This observation led to the Working-Set Model

Working-Set Model (Cont.)

- During the lifetime of the process, references are confined to a subset of pages

Working-Set Model (Cont.)

- $\Delta \equiv$ *working-set window* $\equiv$ a fixed number of page references

  For example: if $\Delta = 10$ memory references, then the working set at t_1 is $\{1,2,5,6,7\}$, at t_2 is $\{3,4\}$

page reference table

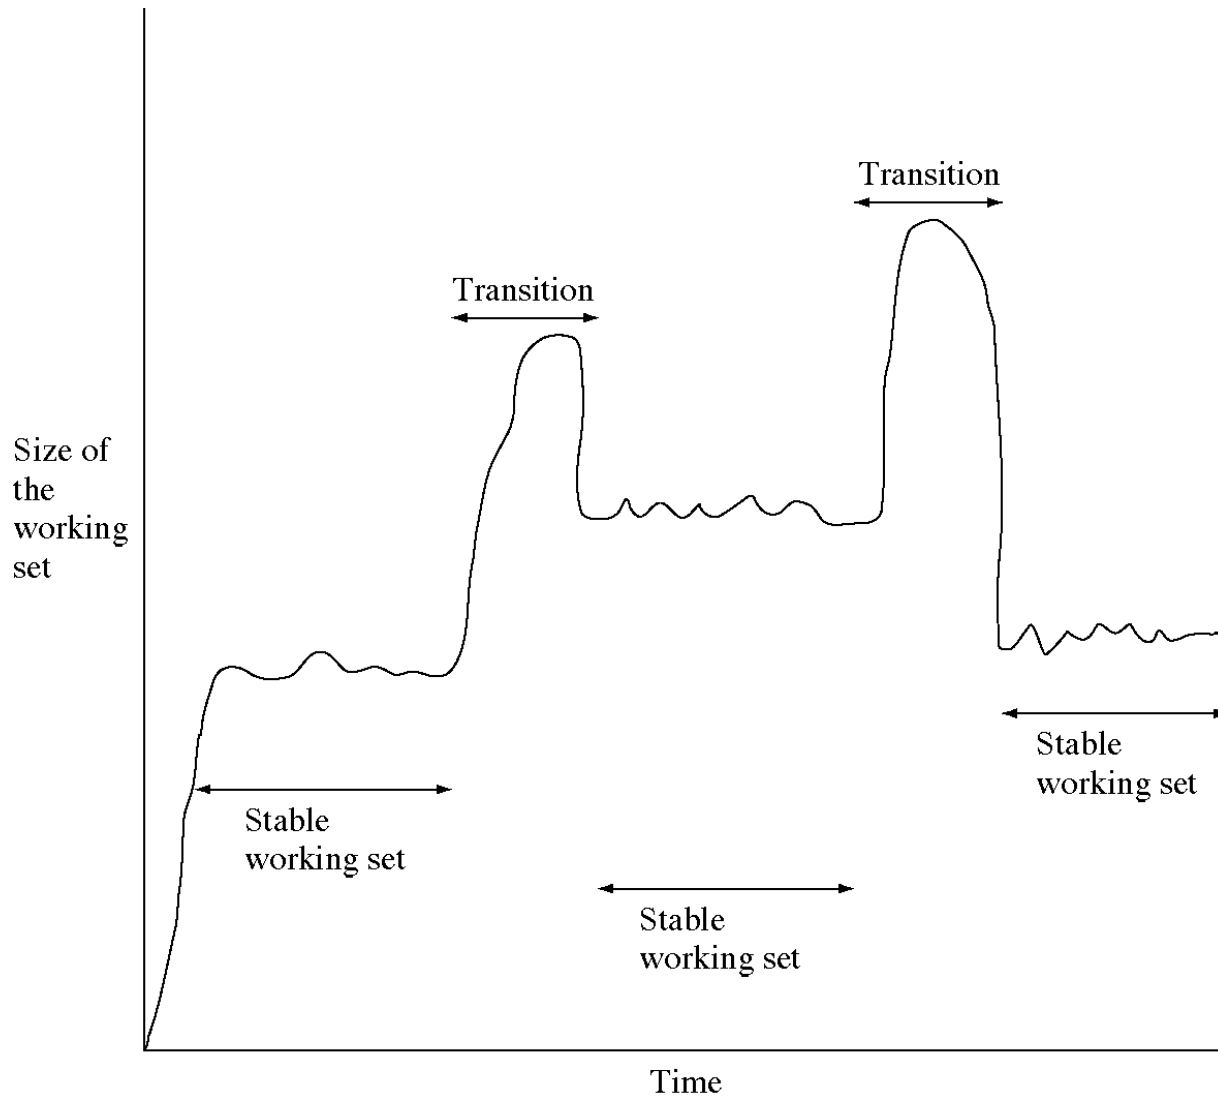. . . 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 . . .

Δ t_1 Δ t_2

$WS(t_1) = \{1,2,5,6,7\}$ $WS(t_2) = \{3,4\}$

Working-Set Model (Cont.)

Size of the working set

Transition

Transition

Stable working set

Stable working set

Stable working set

Time

7.44

Virtual Memory

Working-Set Model (Cont.)

- WSS_i (working set size of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)

- $D = \Sigma \; WSS_i \equiv$ total demand for frames

  m = total # of frames

- if $D > m \; \Rightarrow$ Thrashing

- Policy: if $D > m$, then suspend (i.e., swap out) one of the processes

Page-Fault Frequency Scheme

Page-Fault Frequency Scheme (Cont.)

- Establish "acceptable" page-fault rate
  - If actual rate of a process is too low, remove a frame from that process
  - If actual rate too high, give that process a frame

- May have to suspend (i.e., swap out) a process if the page fault rate increases and no free frames are available

Summary

- Virtual Memory
  - To allow a large logical address space to be mapped onto a smaller physical memory
  - Demand paging is a way to implement virtual memory

- Page Replacement Algorithms
  - FIFO, Optimal, LRU, LRU Approximation

- Frame Allocation
  - Fixed vs. Dynamic Allocation
  - Local vs. Global Allocation

Summary (Cont.)

- Thrashing: process is busy bringing pages in and out (no work is being done)
  - Working-set Model
  - Page-Fault Frequency Scheme